

Technical University of Applied Sciences Würzburg-Schweinfurt (THWS)
Faculty of Computer Science and Business Information Systems

Master Thesis

Comparative Analysis of SVG Deep Generative Models

submitted for the completion of degree

**Master of Science
in
Artificial Intelligence**

Fikrat Mutallimov

Submitted on: April 14, 2026

Supervisor: Prof. Dr. Magda Gregorova
Second Examiner: Prof. Dr. -Ing. Erik Schaffernicht

Abstract (en)

Current generative models can produce visually high-quality raster images from input text or image, but the field falls short in vector graphics format due to its complex nature. An SVG file is a document that contains paths, layers, and hierarchy and it means its purpose is not just to represent something semantically, it should be structurally clean and neat. Those requirements create far more problems for building a model that generates SVG. This thesis analyzes SVG generation methods around thirteen recent models to understand how internal design choices affect the structure of output. The analysis is shaped around three dimensions, how models represent vector geometry internally, what supervision signals guide the generation process and which metrics are useful to evaluate the SVG. Representation and supervision interact rather than operate independently and they are two main parts of internal design of models. The analysis showed that the same loss function can produce different structural outcomes depending on the representation choice and the same representation behaves differently under different supervision. The most consistent pattern across the models is a trade-off between visual realism and editability. Models that generate the most visually detailed outputs also produce the most fragmented SVG files, while methods that constrain their representation or learn from data produce cleaner results, but within narrower domains. Current evaluation metrics (FID, CLIP, LPIPS) all operate on rendered pixels and cannot detect structural differences. To address this, the thesis proposes a six-dimension evaluation framework that includes path economy, geometric quality, layer organization, and editability alongside visual metrics. The thesis does not propose a new model. It provides a framework for understanding why current models produce the structures they do and what needs to change for generated SVGs to work in real design tools.

Abstract (de)

Aktuelle generative Modelle können visuell hochwertige Rasterbilder aus Text oder Bildeingaben erzeugen, doch im Bereich der Vektorgrafiken bleibt das Feld aufgrund deren komplexer Natur zurück. Eine SVG-Datei ist ein Dokument, das Pfade, Ebenen und Hierarchie enthält, und das bedeutet, ihr Zweck ist nicht nur, etwas semantisch darzustellen, sondern sie muss auch strukturell sauber und ordentlich sein. Diese Anforderungen schaffen weit mehr Probleme, die es beim Bau eines Modells zur SVG-Generierung zu überwinden gilt. Die vorliegende Arbeit analysiert SVG-Generierungsmethoden anhand von dreizehn aktuellen Modellen, um zu verstehen, wie interne Entwurfsentscheidungen die Struktur der Ausgabe beeinflussen. Die Analyse ist um drei Dimensionen aufgebaut: wie Modelle die Geometrie von Vektorgrafiken intern repräsentieren, welche Supervisionssignale den Generierungsprozess leiten und welche Metriken zur Bewertung der SVG nützlich sind. Repräsentation und Supervision wirken zusammen statt unabhängig

voneinander und sind die zwei Hauptbestandteile des internen Designs der Modelle. Die Analyse hat gezeigt, dass dieselbe Verlustfunktion unterschiedliche strukturelle Ergebnisse erzeugen kann, je nach Repräsentationswahl, und dieselbe Repräsentation sich unter verschiedener Supervision unterschiedlich verhält. Das konsistenteste Muster über die Modelle hinweg ist ein Kompromiss zwischen visuellem Realismus und Editierbarkeit. Modelle, die visuell besonders detaillierte Ausgaben erzeugen, produzieren gleichzeitig die fragmentiertesten SVG-Dateien, während Methoden, die ihre Repräsentation einschränken oder aus Daten lernen, sauberere Ergebnisse liefern, jedoch innerhalb engerer Domänen. Aktuelle Evaluierungsmetriken (FID, CLIP, LPIPS) arbeiten ausschließlich auf gerenderten Pixeln und können strukturelle Unterschiede nicht erkennen. Um dies zu adressieren, schlägt die Arbeit ein sechsdimensionales Evaluierungsframework vor, das Pfadökonomie, geometrische Qualität, Ebenenorganisation und Editierbarkeit neben visuellen Metriken einbezieht. Die Arbeit schlägt kein neues Modell vor. Sie liefert einen Rahmen, um zu verstehen, warum aktuelle Modelle die Strukturen erzeugen, die sie erzeugen, und was sich ändern muss, damit generierte SVGs in realen Designwerkzeugen funktionieren.

Acknowledgment

I want to thank Prof. Dr. Magda Gregorova and Philipp Vaeth for supervising this thesis. At the beginning of the process, I had no clear direction and their guidance is what kept me focused and moving forward. Every session gave me something concrete to work with. I am also grateful for the freedom they gave me to shape the analytical approach on my own terms.

I thank Prof. Dr. Erik Schaffernicht, as second examiner.

I thank Shabnam Huseynova for her support during the thesis period.

Special thanks to Ismail Eyyublu, a designer at McKinsey in Baku, who helped me understand SVG from the perspective of someone who actually works with it. His explanations of what designers expect from generated vector graphics shaped how I thought about editability throughout this thesis.

AI language model Claude (Anthropic) [1] was used during the preparation of this thesis for brainstorming, literature comprehension, text editing, and translation.

Contents

1	Introduction	1
1.1	Motivation	1
1.2	Research Problem	2
1.3	Thesis Objective	3
1.4	Research Goals	3
1.5	Research Questions	4
1.6	Thesis Structure	4
2	Background	6
2.1	Vector Graphics Fundamentals	6
2.1.1	Paths and Bézier Curves	7
2.1.2	Coordinate Systems and Transformations	8
2.2	Vector Representations	8
2.2.1	Parametric Representation	9
2.2.2	Sequential Representation	9
2.2.3	Implicit Neural Representation	9
2.2.4	Latent Space as a Generative Layer	10
2.3	Differentiable Rendering	10
2.4	Loss Functions and Semantic Guidance	12
2.4.1	Reconstruction Losses	12
2.4.2	CLIP	13
2.4.3	Diffusion Models	13
2.4.4	Score Distillation Sampling and VSD	13
3	Architectural Landscape of SVG Generation	15
3.1	Optimization-Based Methods	16
3.1.1	Reconstruction-Supervised	17
3.1.2	Semantic-Supervised	17
3.1.3	Neural Implicit Representation	19
3.2	Dataset-Driven Methods	20
4	Vector Representation in SVG Generation	23
4.1	Parametric Representations	23
4.2	Sequential Representations	24
4.3	Implicit Neural Representations	24

4.4	Latent Space as a Generative Layer	25
4.5	Representation Boundaries in Practice	25
4.6	Representation, Editability, and Structural Outcomes	26
4.7	Summary	28
5	Supervision and Guidance Strategies for SVG Generation	29
5.1	Reconstruction-Based Supervision	29
5.2	Semantic Guidance with Vision–Language Models	30
5.3	Diffusion-Based Supervision	31
5.4	Dataset Supervision vs. Per-Instance Supervision	32
5.5	Initialization and Structural Priors	32
5.6	Summary	34
6	Evaluation of SVG Generation	35
6.1	Evaluation Metrics in Current Practice	35
6.2	Evaluation Gaps	37
6.2.1	Structural Redundancy	37
6.2.2	Geometric Quality	37
6.2.3	Editability	38
6.3	Efficiency and Reproducibility	38
6.4	Toward a Multi-Dimensional Evaluation Framework	41
6.4.1	Semantic Alignment	41
6.4.2	Path Economy	41
6.4.3	Geometric Quality	41
6.4.4	Layer Organization	41
6.4.5	Editability	42
6.4.6	Efficiency	42
6.4.7	Applying the Framework	42
7	Comparative Synthesis	44
7.1	How Representation and Supervision Shape Structure	44
7.2	The Realism–Editability Trade-off	45
7.3	Conditions for Clean and Editable SVG	47
8	Conclusion	48
8.1	Answers to the Research Questions	48
8.2	Contributions	50
8.3	Limitations	50
8.4	Future Work	51
8.5	Personal Takeaways	51
	Bibliography	53
	List of Figures	57

A	Exclusion of Large Language Model Approaches	60
B	Model Overview Tables	62
B.1	Task and Inference Overview	63
B.2	Training Architecture: Optimization-Based Methods	64
B.3	Training Architecture: Dataset-Driven Methods	65
B.4	Reproducibility Overview	66
C	Architectural Pipeline Diagrams	67
	Declaration on oath	71
	Consent to plagiarism check	72

1 Introduction

Vector graphics differ fundamentally from raster images. Raster images represent visual content as grids of pixels, where each pixel stores color information independently of its neighbors. Vector graphics instead describe images using mathematical primitives: lines, curves, shapes, and groups defined parametrically. Because the geometry is encoded, rather than sampled, vector graphics scale without loss and can be edited at the level of individual components. Among vector formats, Scalable Vector Graphics (SVG) has become the dominant web standard. Its structured nature makes it programmable: SVG files can be generated, modified, and animated through code, which is why the format underlies most of what is drawn on the web. A 2024 survey by W3Techs found that over 65% of all websites use SVG for their graphics [2].

In terms of generative modelling, the structure of SVG introduces challenges that do not arise in raster synthesis. SVG graphics encode both geometry and topology. So a generated SVG must satisfy two requirements simultaneously: it must look correct and it must be structurally clean. Although it may seem visually plausible, an output with redundant paths and unstable geometry cannot be edited or reused. Visual quality and structural quality are distinct properties, and current evaluation practice does not consistently distinguish between them.

Before learning-based approaches, vector graphics were produced through algorithmic vectorization: tools that approximate raster images using geometric primitives via edge detection, contour tracing, and polygon approximation. These methods remain in use but rely on heuristic rules and struggle with complex scenes. Learning-based approaches introduce new possibilities, but also raise questions about how models construct vector structure internally and whether the outputs they produce are usable in practice.

1.1 Motivation

In recent years, generative models have substantially changed what is possible in visual content creation. Diffusion models, large-scale multimodal systems, and transformer based architectures have demonstrated that high-quality image synthesis from text prompts is no longer a research prototype. Systems such as Stable Diffusion, DALL-E,

and Midjourney produce raster images that are often indistinguishable from photographs or hand-drawn illustrations. Progress in this area has been fast. But this progress has not extended equally to vector graphics. Research on SVG generation has grown, but the reliability and structural quality of generated outputs remain open problems. It is worth considering that vector graphics have not become less important. Instead, icons, logos, interface elements, illustrations, and data visualizations are created and maintained in vector formats precisely because they remain editable and resolution-independent. The formats used by Adobe Illustrator, Figma, and every major web browser are vector-based.

The gap between raster and vector generation is not only technical difficulty. It reflects a mismatch between how generative models are trained and evaluated and what vector graphics actually require. A model trained to minimize pixel reconstruction error has no reason to produce clean paths or interpretable layers or a model evaluated by FID score is rewarded for visual plausibility, not structural usability. Those are just a few examples, and we will see more of them later. It is understandable, because early generative models focused on raster images which led the vector graphics field to be partially dependent on raster images and their features. But still there are different methods and techniques used to generate SVG. Rather than proposing a new generation model, this thesis analyzes existing methods from the inside: how they represent vector graphics, what supervision signals they use, and how those choices shape the structure of the outputs they generate.

1.2 Research Problem

SVG generation research has divided across different directions. Some models treat SVG as a stream of drawing commands, other approaches refine vector parameters directly against a raster target. It is even possible to define geometry through learned functions.

These methods differ not only in architecture but in how they internally represent vector structure and what signals they use to guide generation. Despite this diversity, a systematic understanding of how these design choices influence the structure of what gets produced is missing. Models use different representations, different supervision signals, and different optimization strategies, but no existing work traces how those choices connect to the structural properties of the output. Two models may achieve similar FID scores while producing outputs with completely different path counts, layer organization, and editability. The field has many methods but no shared framework for understanding why they produce the structures they do.

Evaluation makes this harder. The metrics most commonly used (FID, LPIPS, CLIP score) are all borrowed from raster synthesis and operate on rendered pixels. They

cannot detect redundant paths, broken hierarchy, or unusable geometry. There is no widely adopted benchmark for SVG editability, path economy, or layer coherence. This combination — heterogeneous architectures, untraced design-to-structure relationships, and evaluation that ignores structure — constitutes the research gap addressed by this thesis.

1.3 Thesis Objective

This thesis investigates SVG generation methods from the perspective of representation, supervision, and structural quality. The goal is not to propose a new model but to develop a structured analytical understanding of how existing methods construct vector graphics and why their structural outcomes differ.

The analysis focuses on models where the generation process is interpretable in terms of vector operations or clearly defined architectural components. Large language model approaches, which treat SVG generation as code synthesis, are reviewed as a distinct and growing direction but excluded from the main analysis. A full justification is provided in Appendix A.

The thirteen models in the main analysis were selected by a few practical criteria, and the selection is not exhaustive. Each one is a recent method (2020 to 2025) that uses differentiable rendering or learned models and produces SVG or structured vector output. Each paper also describes its architecture in enough detail to analyze how the model represents geometry and how it is supervised. The models were chosen so that together they cover the different representation and supervision types. Whether code is publicly available was recorded (Table 6.2), but it was not a requirement.

1.4 Research Goals

1. To explain and analyze the core properties, structures, and components of SVG graphics as a foundation for understanding their representation and generation.
2. To collect, review, and categorize a representative set of existing methods for SVG generation across different input modalities.
3. To investigate the deep learning architectures employed in SVG generation models and analyze their design principles.

4. To examine how these architectures represent SVG data internally and how outputs are transformed into standard SVG files.
5. To evaluate existing approaches, metrics, and benchmarks for assessing the quality of generated SVGs.
6. To assess and compare the structural and visual quality of current SVG generation models based on identified evaluation criteria.

1.5 Research Questions

1. What are the main properties and structures that define an editable and high-quality SVG?
2. What methods currently exist for generating SVGs from different input modalities?
3. What types of deep learning architectures are used in SVG generation models?
4. What kinds of structure do these models produce?
5. How can the quality and editability of generated SVG be evaluated?
6. What is the reported and assessed quality of SVG outputs across different models?

These questions are not independent. Together they form a single analytical inquiry: how do the internal design choices of SVG generation models (what they represent, how they are supervised, and how their outputs are evaluated) determine whether the resulting graphics are structurally usable as well as visually plausible. The thesis argues that this question cannot be answered by visual metrics alone, and that systematic analysis of representation and supervision is necessary to explain the structural differences observed across current methods.

1.6 Thesis Structure

Chapter 2 introduces the foundational concepts required to understand vector graphics generation: SVG structure, representation types, differentiable rendering, and the supervision strategies used in learning-based methods. Building on this, Chapter 3 surveys

existing SVG generation approaches and categorizes them by architectural family and input modality. Chapters 4 and 5 analyze how representation and supervision choices shape vector structure, tracing the relationship between internal design decisions and observable output properties. Chapter 6 introduces a multi-dimensional evaluation framework for SVG generation that extends beyond visual similarity to include structural and practical criteria. These perspectives are then applied in Chapter 7, where strengths, limitations, and trade-offs across the analyzed models are examined comparatively. The thesis concludes by mapping findings to the research questions, summarizing contributions, and identifying directions for future work.

2 Background

Generating vector graphics computationally is harder than it seems. Unlike raster images, which can be produced by predicting pixel values, SVGs have to be structurally correct and the resulting file needs to be something a designer can actually open and work with. Otherwise, the output is not really something that can be used in practice. Getting there requires understanding the building blocks that most SVG generation models have in common, despite the variety of approaches and architectures that exist. This chapter covers those building blocks. It starts with SVG itself: what the format represents, how it is structured, and why those structural properties create constraints that raster generation does not have. From there it introduces four internal representation types used across the analyzed models, then differentiable rendering, which is the technical bridge that connects vector parameters to pixel-space losses. The chapter closes with the supervision signals (reconstruction losses, CLIP, SDS, VSD) that drive the generation process in different ways.

2.1 Vector Graphics Fundamentals

Scalable Vector Graphics (SVG) is an XML-based standard for representing two-dimensional graphics as structured text [3]. The difference between raster formats and SVG is that an SVG file is a readable document that describes how an image should be drawn rather than storing pixel values. At its core, an SVG file consists of instructions or commands that define geometric primitives such as circles, rectangles, lines, and paths, along with styling and hierarchical relationships between elements [4].

SVG supports a range of graphic elements, including basic shapes such as lines, rectangles, and circles, as well as text and raster images. The most flexible element is the path, which allows complex shapes to be defined using sequences of commands: moving a virtual pen, drawing straight lines, and creating smooth curves [3].

What makes SVG practically useful is editability, and that editability comes from groups and layers. Multiple shapes can be combined into groups that behave as a single unit for styling and editing [4]. Elements are rendered in a defined order, and later elements appear on top of earlier ones, making layering an intrinsic part of the representation.

This hierarchical structure is what makes SVG files editable in practice: a designer can select a group, move a layer, or change one element without affecting the rest.

Figure 2.1 shows a minimal SVG file and its rendered output. The file is plain text: each element is a single readable line that defines geometry and color directly. This is what makes SVG editable. A designer can open the file, find the element they want, and change a number.

```
<svg xmlns="http://www.w3.org/2000/svg"
      width="200" height="200">
  <!-- Face -->
  <circle cx="100" cy="100" r="80"
          fill="#FFD700"/>
  <!-- Left eye -->
  <circle cx="70" cy="80" r="10"
          fill="#333"/>
  <!-- Right eye -->
  <circle cx="130" cy="80" r="10"
          fill="#333"/>
  <!-- Sad mouth -->
  <path d="M 65 135 Q 100 110 135 135"
        stroke="#333" stroke-width="5"
        fill="none"/>
</svg>
```



Figure 2.1: A minimal SVG file (left) and its rendered output (right). Each element is one readable line: the face, eyes, and mouth are independent objects defined by plain text parameters. Any of them can be selected and changed without touching the rest.

2.1.1 Paths and Bézier Curves

Most SVG graphics are built using paths. A path is a sequence of commands (*move to*, *line to*, *curve to*) that define its shape [3]. Curved segments use Bézier curves, which are not defined by the points they pass through but by control points that pull the curve toward them from the outside. Moving a control point reshapes the curve smoothly without breaking it.

As shown in Figure 2.2, SVG supports two types. Quadratic Bézier curves use a single control point and are simpler but less flexible. Cubic Bézier curves use two control points and can represent almost any smooth shape [3]. Most generation models in this thesis work with cubic curves because they are expressive enough to approximate complex

geometry. The control point coordinates are the core parameters that models either predict directly or adjust during optimization.




Figure 2.2: Bézier curves of degree 1, 2, and 3 (red) with their control polygons (gray). From left to right, an additional control point (blue) is added at each degree. The curve changes direction and curvature as control points are inserted or moved. Image: Wikimedia Commons [5].

2.1.2 Coordinate Systems and Transformations

SVG uses a simple coordinate system. The top-left corner of the canvas is the origin, x increases to the right, and y increases downward [4]. Every element is positioned within this grid using numeric coordinates, the same numbers you see in the SVG file.

What makes this system powerful for design is transformations. Any shape or group can be moved, scaled, rotated, or skewed by applying a transform attribute. When a transformation is applied to a group, every element inside moves together. Groups can be nested inside other groups, each with their own local coordinate frame. A logo might have a top-level group for the whole graphic, sub-groups for each component, and individual paths inside those. Moving the top-level group moves everything. Moving a sub-group moves only that part. This nested structure is what allows complex SVG files to remain organized and editable even as they grow in size.

2.2 Vector Representations

SVG generation and vectorization models do not all manipulate vector graphics the same way internally in generation processes. Obviously, the final output is always a text-based file or SVG itself as image. What differs is the internal representation used during learning or optimization. This choice directly affects which algorithms can be applied,

how stable training is, what structural properties tend to emerge and the quality of final output. The following sections introduce the four representation types that appear across the models analyzed in this thesis.

2.2.1 Parametric Representation

Parametric representation is natural for SVG because SVG is already parametric. Every element in the file is defined by control point coordinates, stroke widths and colors which are readable and editable as plain text [3]. The introduction of differentiable rasterization by Li et al. [6] made it practically possible to optimize these parameters directly using image-based loss functions: a graphic is initialized as a collection of primitives, and their parameters are iteratively refined until the rendered output matches a target.

2.2.2 Sequential Representation

The idea of treating vector graphics as ordered sequences of drawing commands goes back to SketchRNN by Ha and Eck [7], who showed that recurrent networks could learn to generate stroke-based drawings by modelling them as sequences of pen movements. Vector graphics encoded this way look like a script: move the pen here, draw a line there, curve to this point. These representations share the same basic structure as a sentence, which opened the door to applying language modelling techniques directly. The downside of that approach is that whenever errors accumulate, a mistake early in a long sequence can break everything that follows.

2.2.3 Implicit Neural Representation

The use of neural networks to define geometry implicitly through learned weights was established in 3D scene representation by Mildenhall et al. [8]. In NeRF, the network acts like a function: based on a given location, it tells you what occupies that space. Nothing is stored explicitly, the geometry only appears when you query the network. Thamizharasan et al. [9] adapted this idea to vector graphics by replacing explicit path coordinates with neural implicit layers that define shape boundaries indirectly. The network learns to represent the geometry as a continuous function and the resulting shapes emerge from the learned implicit representation instead of being directly parameterized. This allows for more flexible and compact representations, but also makes it harder to control specific geometric features directly, because they are not explicitly defined as parameters.

Take the mouth in Figure 2.1, `M 65 135 Q 100 110 135 135`. A parametric representation stores its three points directly: start (65, 135), control point (100, 110), and end (135, 135), plus stroke width and color. Optimization changes these numbers. A sequential representation turns the same curve into tokens: a move command with the arguments 65 and 135, then a quadratic curve command with 100, 110, 135, and 135. The model has to predict them in that order. An implicit representation like NeuralSVG stores none of these numbers. The mouth is simply shape number 4. The MLP takes the index 4 as input and outputs the control points and the color of that shape [10]. To change the mouth, you change the network weights, and the numbers only appear when the network is queried. NIVeL works the other way around: its network takes a pixel coordinate and returns which layer covers it, so the curve only appears after it is traced from the field [9].

2.2.4 Latent Space as a Generative Layer

Latent representation is not a geometric representation at all. It is a generative layer that sits on top of one of the other three. There are no control points, no path commands, no shape boundaries stored anywhere. Latent space compresses many graphics into a shared numerical embedding learned during training. Kingma and Welling [11] formalized this idea through the Variational Autoencoder: an encoder compresses an input into a short vector and a decoder reconstructs the original from it.

For SVG generation, the actual geometry is always produced by the decoder, which translates a latent vector back into one of the other representation types: parametric paths or sequential commands. The latent space sits on top as an organizational layer. It is included here because several models use it this way and confusing it with a geometric representation leads to misunderstanding how those models actually construct their outputs.

2.3 Differentiable Rendering

DiffVG [6] is one of the main bridges between vector and pixel spaces. Before neural networks can optimize vector graphics using image-based loss functions, there needs to be a way to connect the two spaces. Vector parameters (control points, stroke widths, colors) live in a continuous geometric space. Loss functions operate on pixel values. These two spaces are not naturally connected: traditional rasterization converts vector geometry into pixels, but this process is not differentiable, meaning gradients cannot flow back through it to update the vector parameters.

DiffVG, introduced by Li et al. [6], solves this by approximating rasterization in a way that supports gradient computation. Given a set of vector parameters as input, DiffVG renders a pixel image as output. Because the rendering step is differentiable, the loss computed on the output image (whether a pixel reconstruction loss, a perceptual loss, or a signal from a pretrained semantic model) can be backpropagated all the way through to the vector parameters themselves. The system learns which control point to move, and in which direction, to reduce the error. The tricky part is handling edges and anti-aliasing, places where a tiny shift in a path boundary causes a sudden pixel change. DiffVG approximates these discontinuities in a way that still produces usable gradients [6].

Before DiffVG, optimizing vector graphics with gradient-based methods was not straightforward. Its introduction made it possible to apply the full range of image-based supervision strategies, reconstruction losses, CLIP guidance, diffusion priors directly to vector parameters.

The core idea is shown in Figure 2.3. Normally, converting vector shapes to pixels is a one-way street: you can go from shapes to pixels, but not back. DiffVG makes the middle step reversible. Once gradients can flow back through the rasterizer, any image-based loss function becomes usable for optimizing vector parameters directly.




Figure 2.3: The differentiable rendering pipeline. DiffVG converts vector parameters into a pixel image. A loss compares this image to a target. Gradients then flow back through DiffVG to update the vector parameters directly [6].

Not every analyzed model needs differentiable rendering. DeepSVG and DeepIcon are trained directly on SVG commands, with losses that compare predicted and ground-truth commands, so no rasterizer is involved [12, 13]. LayerTracer generates raster construction frames with a diffusion model and vectorizes them afterwards with a separate tracing step [14]. NIVeL optimizes implicit fields on a pixel grid and extracts curves only at the end [9]. DiffVG is also not the only differentiable rasterizer. Bézier Splatting, for example, represents each curve as 2D Gaussians placed along it and reports much faster forward and backward passes than DiffVG [15].

2.4 Loss Functions and Semantic Guidance

The supervision signals in this section fall into two groups. Reconstruction supervision compares the rendered output with a target image (Section 2.4.1). Semantic supervision has no target image and uses a pretrained model, CLIP or a diffusion model, to judge whether the output matches a text prompt (Sections 2.4.2 to 2.4.4). Chapter 3 uses this split to group the optimization-based methods.

Even when a model manipulates vector parameters internally, comparing outputs in pixel space is more practical than comparing vector structures directly. Most visual data exists as raster images, and pixel-based comparison gives access to well-established loss functions from computer vision. When combined with differentiable rendering, gradients from these losses flow back through the rasterizer to update vector parameters. Which loss you use matters a lot. It shapes what the model focuses on during training.

2.4.1 Reconstruction Losses

The simplest losses operate directly on pixel values. The L2 loss measures the mean squared difference between predicted and target pixels, while the L1 loss measures the mean absolute difference [16]. Both are computationally efficient and provide stable gradients when a reference image is available. The problem is that they treat every pixel the same. A pixel at a shape boundary matters more than one in a flat colored region, but L1 and L2 do not know that.

Perceptual losses address this by comparing feature representations extracted from pretrained convolutional networks rather than raw pixel values. The formulation proposed by Johnson et al. [17] uses intermediate feature maps from VGG [18], which capture edges, textures, and object-level structure. Comparing at the feature level produces supervision signals that better reflect visual similarity as humans perceive it, and helps preserve overall structure even when exact pixel alignment differs slightly.

For tasks where shape boundary accuracy matters more than color fidelity, overlap-based losses are used. Intersection over Union (IoU) [19] measures the overlap between predicted and target regions. In vector graphics, mask-based supervision of this kind encourages correct topology and spatial consistency across layers, rather than minimizing pixel intensity differences.

All three loss types depend on having a reference raster image. In text-to-SVG settings, where no target image exists, alternative guidance strategies are needed.

2.4.2 CLIP

CLIP (Contrastive Language–Image Pretraining), introduced by Radford et al. [20], is trained on large-scale image-text pairs to learn a shared embedding space where semantically matching pairs are placed close together and mismatched pairs are pushed apart. Given an image and a text prompt, CLIP measures their similarity through separate image and text encoders, typically using cosine similarity between the resulting embeddings.

In SVG generation pipelines, the rendered vector graphic is passed through the CLIP image encoder while the text prompt is encoded separately. If the rendered output does not match the prompt semantically, the similarity score will be low. Because CLIP is differentiable with respect to its image input, and because images can be produced from vector parameters via differentiable rendering, the CLIP similarity score can serve as an optimization signal. CLIP does not generate images in this setup: it acts as a semantic evaluator that provides a direction for improvement. In the models analyzed in Chapter 5, CLIP plays a smaller role. VectorFusion, for example, uses CLIP only to pick the best of several raster samples before vectorization, and the optimization itself is driven by a diffusion model [21].

2.4.3 Diffusion Models

Diffusion models are generative models that learn to reverse a gradual noise corruption process. During training, clean images are progressively corrupted by adding noise, and the model learns to denoise them step by step. After training, new images can be generated by starting from random noise and iteratively denoising [22]. Stable Diffusion [23] extends this by performing the diffusion process in a compressed latent space rather than directly in pixel space, significantly reducing computational cost while preserving image quality. When conditioned on text, diffusion models learn to generate images consistent with the textual description, capturing complex visual distributions and rich semantic structure.

2.4.4 Score Distillation Sampling and VSD

Score Distillation Sampling (SDS), introduced by DreamFusion [24] for text-to-3D generation and later adopted for vector graphics by VectorFusion [21], uses a pretrained text-conditioned diffusion model as an optimization signal. The current SVG is rendered to an image, noise is added according to the diffusion schedule, and the diffusion model predicts which noise was added, conditioned on the text prompt. The difference between

the predicted noise and the noise that was actually added, weighted according to the timestep, defines a gradient, which is backpropagated through the rasterizer to update the SVG parameters. No ground-truth image is required. The diffusion model’s learned prior acts as the teacher.

A known limitation of SDS is that results look oversaturated and over-smoothed, with little fine detail. SDS only works well with a very high guidance weight, and that weight pushes colors to extremes. It also pulls every run toward the same most likely image, so detail and variety get lost [24, 25].

Variational Score Distillation (VSD) was introduced by ProlificDreamer [25] for text-to-3D. Instead of one output, it keeps a few versions at once, called particles. A LoRA copy of the diffusion model (a small set of extra low-rank weights added to the frozen model) is trained next to them and learns what they currently look like. This works with normal guidance weights and gives more detail and variety. SVGDreamer builds on it with Vectorized Particle-based Score Distillation (VPSD), where each particle is a set of SVG parameters [26]. Both SDS and VSD are used across several models analyzed in Chapter 5. The complete optimization loops for SDS and VSD are shown in Figures 3.2 and 3.3, where the role of the diffusion model as a supervision signal becomes concrete. Chapter 3 now introduces the thirteen models that are built from these blocks.

3 Architectural Landscape of SVG Generation

SVG generation is not a single problem with a single solution. Some methods treat vector graphics as optimization targets and adjust geometric parameters until a rendered image matches a prompt or a reference image. Others learn from datasets and generate SVGs through a trained decoder. A group of models avoids explicit geometric primitives entirely and uses neural functions to define shapes indirectly. These approaches reflect fundamentally different assumptions about what a good SVG output should look like, and they lead to very different trade-offs between visual quality and structural organization.

This chapter introduces the thirteen models that form the whole thesis. They were selected among a broader research landscape that also includes sequence-based sketch models, retrieval-based vectorization systems, and large language model approaches to SVG synthesis. That wider context is sketched below to clarify where the selected models sit and why they were chosen. The organizing principle is simple: these thirteen models together cover the main design decisions of models in two dimensions already mentioned: representation and supervision, that shape how SVG structure emerges from a generation pipeline. A summary of all thirteen models is provided in Table 3.3, and Figure 3.1 shows how they are grouped in this thesis.

An early line of work treats vector drawings as ordered sequences of commands or strokes. Starting from SketchRNN (Section 2.2.2), later methods extended this to full SVG command prediction, including Sketchformer [27] and related sequence models [28, 29]. These approaches established that SVG structure could be learned statistically, but long command chains make training difficult and generalization to complex scenes remains limited by the scarcity of large-scale SVG datasets.

A second direction introduced differentiable rasterization as a bridge between vector parameters and image-space supervision. This opened the door to reconstruction-based vectorization and, later, to semantically supervised generation using pretrained diffusion models. Most of the optimization-based models in this thesis belong to this lineage.

A third direction replaces explicit geometric primitives with neural implicit functions. Instead of optimizing Bézier coordinates directly, these methods run optimization in neural parameter space, where an MLP defines shape boundaries indirectly. This approach has




Figure 3.1: Taxonomy of the thirteen analyzed models by generation strategy.

connections to neural radiance fields [8] and represents an attempt to improve structural coherence through representational design rather than supervision alone.

More recently, large language models and multimodal variants have entered SVG generation. Methods such as StarVector [30], OmniSVG [31], LLM4SVG [32], and InternSVG [33] treat SVG as structured code and generate it through transformer architectures pretrained on a collection of large-scale texts and images. In one sense, SVG is a natural fit for language models: the file format is sequential XML commands, not unlike a programming language. Visual alignment improves substantially compared to earlier sequence models, and these approaches handle a wide range of generation and editing tasks. However, they are not included in the main analysis here. The core question of this thesis is how specific architectural decisions impact the structure of the output and that relationship is difficult to trace through billions of parameters. A more detailed justification is provided in Appendix A.

3.1 Optimization-Based Methods

Optimization-based methods share a common logic: SVG primitives are randomly or heuristically initialized and then iteratively refined until the rendered output satisfies a supervision signal which is either prompt or target image. The SVG is the variable being optimized, not the output of a forward pass. What varies across methods in this family is what that supervision signal is and how geometry is represented during optimization.

Almost all of them rely on DiffVG [6] to work. NIVeL is the exception: it optimizes

implicit fields that are evaluated on a pixel grid and only turned into curves afterwards with marching squares, so DiffVG is not part of its optimization loop [9]. The optimization loop across all methods follows the same render-compare-update cycle, but differs in what loss drives it and how geometry is represented. Table 3.1 summarizes how each model in this group implements those differently. Fuller versions of Tables 3.1 to 3.3 are in Appendix B.

3.1.1 Reconstruction-Supervised

LIVE [34] is the simplest form of optimization-based vectorization. It starts from an empty canvas, finds where the biggest visual gap is between the current SVG and the target image, drops a new path there, and optimizes everything together. It repeats this process layer by layer until the reconstruction is close enough. Its loss combines a pixel loss weighted toward shape contours (UDF) with a Xing loss that penalizes self-intersecting paths [34]. Because paths are added one at a time from coarse to fine, the output tends to have a layered structure where each path has a recognizable role, at least at small path counts. With many paths this breaks down (Section 4.1). SAMVG [35] follows the same optimization loop but starts from a better place: instead of finding gaps by pixel difference, it uses the Segment Anything Model [36] to identify semantically meaningful regions before any optimization begins. The loss is the same kind, pixel reconstruction. Much of the difference in output quality comes from where optimization starts (Section 5.5). Both pipelines are shown in Figure C.1.

Im2Vec [37] does not fit cleanly into either group and is listed as a hybrid in Figure 3.1. It trains an encoder-decoder to predict vector primitives from a raster image in a single forward pass, using differentiable rendering and reconstruction loss during training. The training data are raster images only, the model never sees an SVG [37]. At inference there is no iterative refinement. It is included here because its core contribution is the training objective, not a learned generative model over SVG structure.

3.1.2 Semantic-Supervised

When no target image exists, the optimization loop needs a different kind of signal. These methods use a pretrained diffusion model to judge whether the current SVG looks like what the text prompt describes. The mechanism behind this is Score Distillation Sampling (SDS) [24], originally developed for text-to-3D generation and described in Section 2.4.4: render the SVG, add noise, let the diffusion model predict that noise and use the difference between predicted and added noise as a gradient to update the SVG parameters.

VectorFusion [21] was among the first to apply SDS to SVG generation, using Stable Diffusion [23] as the supervision backbone (Figure 3.2). The outputs are semantically aligned with the prompt but tend to be over-smoothed, a known side effect of standard SDS. SVGDreamer [26] addresses this by replacing SDS with Vectorized Particle-based Score Distillation (VPSD, Section 2.4.4), which models the SVG as a distribution of particles rather than a single point estimate (Figure 3.3). SVGDreamer also introduces explicit control over how many primitives are used and how they are organized into semantic layers, through its semantic-driven image vectorization (SIVE). SVGDreamer++ [38] extends this further with hierarchical image vectorization (HIVE), which works on objects and their parts, and with an adaptive control that can add or remove primitives during optimization. The pipelines of all three are compared in Figure C.2.




Figure 3.2: SDS optimization loop in VectorFusion [21].




Figure 3.3: VPSD optimization loop in SVGDreamer [26].

The two loops look alike, but they differ in what is trained. In SDS (Figure 3.2) everything except the SVG is frozen. The render is encoded and noised, the frozen model predicts the noise, and the difference between predicted and added noise flows back through the

DiffVG into one SVG. VPSD (Figure 3.3) keeps the diffusion model frozen too, but adds two things.

There are k SVGs (particles) instead of one. A LoRA copy of the diffusion model is trained on their renders, and the SVG update comes from the difference between the frozen model and this copy. A reward model also scores the particles and reweights them. The authors report this gives better-looking results and faster convergence [26].

3.1.3 Neural Implicit Representation

NIVeL [9] and NeuralSVG [10] take a different approach to the same optimization loop. They store geometry as the weights of a small MLP, instead of storing geometry as explicit Bézier coordinates that are adjusted directly. In NeuralSVG the network takes a shape index and outputs the geometry for that shape. In NIVeL it takes a pixel coordinate and returns which layer covers it. Nothing is stored as an explicit path.

This brings different properties to optimization-based approaches. In parametric methods each path is an independent object, and optimization can push the paths in different directions. With an MLP, all shapes come from the same learned function, which creates a bias toward global coherence. NIVeL uses this to decompose scenes into layered components, where each layer is a separate neural field whose zero-level set defines a shape boundary. NeuralSVG assigns each shape a fixed index (Section 4.3). The outputs from both tend to be cleaner and more organized than purely parametric optimization, but they still need a full optimization run for every new prompt. Figure C.3 compares the two pipelines.

Table 3.1: Architectural overview of optimization-based SVG generation methods. Trainables: what is updated, per-instance parameters optimized at inference time, except for Im2Vec, whose network weights are trained once. Init.: how primitives are placed before optimization. Supervision: the loss that drives the update. Representation: how geometry is stored while it is optimized.

Model	Trainables	Init.	Supervision	Representation
LIVE	SVG param.	Error-driven	Pixel recon.	Parametric
SAMVG	SVG param.	Mask-based	Pixel recon.	Parametric
Im2Vec	Encoder + RNN	Image encoding	Pixel recon.	Parametric paths
VectorFusion	SVG param.	Random/Trace	SDS	Parametric
SVGDreamer	SVG param. + LoRA	Attention (SIVE)	VPSD	Parametric
SVGDreamer++	SVG param. + LoRA	Random + HIVE	VPSD + HIVE	Structured param.
NeuralSVG	MLP weights + LoRA	Saliency-based	SDS	Neural implicit
NIVeL	MLP weights	Random weights	SDS	Neural implicit

Table 3.1 summarizes the architectural components of the optimization-based methods.

For all models here except Im2Vec, “Trainables” means per-instance parameters optimized at inference time. None of them needs a training dataset. Supervision comes from the input image or from a frozen pretrained model. Im2Vec is different. Its trainables are network weights, learned once from raster images.

3.2 Dataset-Driven Methods

Dataset-driven methods learn structural patterns from collections of vector graphics and generate through a trained forward pass. The optimization happens during training, not at inference. Outputs reflect statistical regularities in the training data rather than the iterative refinement of any single instance. These methods are more constrained by the training corpus but can produce more consistent and visually appealing results without the computational cost of optimization at inference.

DeepSVG [12] is the foundational model in this family. It learns a hierarchical latent space over SVG path commands using a transformer-based encoder-decoder trained on the SVG-Icons8 dataset. The latent space separates path-level and command-level representations, allowing the model to capture both coarse shape organization and fine geometric detail. At inference, all SVG commands are decoded from a sampled latent vector in a single forward pass, not one after another [12]. DeepSVG demonstrated that structured SVG generation from a learned latent space is feasible, and also exposed the difficulty of maintaining geometric validity across long command sequences.

DeepIcon [13] applies a similar framework to icon vectorization specifically, trained on the same SVG-Icons8 dataset of simple icons. The domain-specific training produces outputs with more consistent style and geometry than general-purpose models, though generalization beyond the training distribution is limited by design.

LayerTracer [14] approaches the problem differently. Rather than predicting SVG commands from a latent vector, it models the construction process of vector graphics layer by layer. A diffusion-based architecture generates intermediate structural representations that guide each new layer, approximating the sequential way designers build graphics in practice. This process-aware formulation produces outputs with more meaningful layer organization than direct command prediction, with direct consequences for editability.

T2V-NPR [39] combines dataset training with per-prompt optimization. It first trains a neural path VAE on the FIGR-8-SVG icon dataset, so the latent space learns what a regular, smooth path looks like. For a new prompt there is no forward pass. Path latents in this learned space are optimized with VSD from a pretrained diffusion model, and the result is then refined with layer-wise vectorization [39]. So T2V-NPR is a hybrid: a trained path prior with per-instance optimization on top. It is described in this section

because its path prior comes from a dataset, but its inference works like the optimization methods above.

SVGFusion [40] is the most architecturally complex model in this family. It first learns a continuous latent space of vector graphics using a Vector-Pixel Fusion Variational Autoencoder (VP-VAE), which encodes each SVG together with its rendering. Generation then occurs through a Vector Space Diffusion Transformer (VS-DiT), which operates entirely in this learned latent space. This makes SVGFusion distinct from the semantic-supervised optimization methods in the previous section: where VectorFusion and SVGDreamer use diffusion gradients to guide per-instance optimization in pixel space, SVGFusion uses a diffusion process over a learned vector distribution to sample new SVGs directly. The outputs reflect the statistical structure of the training corpus rather than the convergence of any optimization loop. Table 3.2 summarizes the architectural components of the dataset-driven methods.

Table 3.2: Architectural overview of dataset-driven SVG generation methods. Training data refers to the dataset used for offline model training.

Model	Trainables	Data	Supervision	Representation
DeepSVG	VAE + Transformer	SVG-Icons8	Recon. + KL	Sequential tokens
DeepIcon	Transformer	SVG-Icons8	Reconstruction	Sequential tokens
SVGFusion	VP-VAE + VS-DiT	SVGX	Recon. + Denoising	Continuous latent
T2V-NPR	NPR-VAE + LoRA	FIGR-8-SVG	VSD	Neural path repr.
LayerTracer	DiT + LoRA	20k layered SVGs	Denoising	Layered sequences

Unlike the optimization-based group, every model here has a training corpus. The representation column also shifts away from parametric paths toward sequential tokens, latent spaces, and layered sequences, reflecting a fundamentally different approach to how SVG structure is encoded.

Table 3.3 summarizes the thirteen models analyzed in this thesis. The *Inference* column distinguishes between methods that optimize vector parameters at inference time (*Iterative Opt.*) and methods that generate SVGs through a single trained forward pass (*Forward Dec.*). The *Supervision* column indicates the primary signal driving generation: *Pixel recon.* refers to reconstruction against a target raster image without a pretrained model, *Semantic* to signals derived from frozen pretrained models (diffusion or CLIP), and *Supervised* to training against ground-truth SVG data. *Hybrid* in the Inference column marks methods that combine a trained model with a per-instance optimization or vectorization step. Task types distinguish between *Vectorization* (raster image as input, SVG as output), *Text-to-SVG* (text prompt as input), and *Unconditional Generation* (no input at inference — SVG sampled directly from a learned latent space).

The split between Iterative Opt. and Forward Dec. in the Inference column maps almost cleanly onto the two families: every optimization-based model except Im2Vec refines at

Table 3.3: Overview of the thirteen models analyzed in this thesis, organized by generation strategy.

Model	Task	Inference	Input	Output	Supervision
LIVE	Vectorization	Iterative Opt.	Image	Paths	Pixel recon.
SAMVG	Vectorization	Iterative Opt.	Image	Paths	Pixel recon.
Im2Vec	Vectorization	Forward Dec.	Image	Paths	Pixel recon.
VectorFusion	Text-to-SVG	Iterative Opt.	Text	Paths	Semantic
SVGDreamer	Text-to-SVG	Iterative Opt.	Text	Paths	Semantic
SVGDreamer++	Text-to-SVG	Iterative Opt.	Text	Paths	Semantic
NIVeL	Text-to-SVG	Iterative Opt.	Text	Layers/Paths	Semantic
NeuralSVG	Text-to-SVG	Iterative Opt.	Text	Layers/Paths	Semantic
DeepSVG	Unconditional	Forward Dec.	None	Paths	Supervised
DeepIcon	Vectorization	Forward Dec.	Image	Layers/Paths	Supervised
SVGFusion	Text-to-SVG	Forward Dec.	Text	Layers/Paths	Supervised
T2V-NPR	Text-to-SVG	Hybrid	Text	Layers/Paths	Sup. + Semantic
LayerTracer	Text-to-SVG	Hybrid	Text/Img	Layers/Paths	Supervised

inference time, while every dataset-driven model except T2V-NPR generates through a forward pass. Im2Vec is trained and then decodes in a single pass. T2V-NPR trains a path prior and then optimizes for each prompt. LayerTracer is labeled Hybrid because it combines a trained diffusion model with a per-instance vectorization step.

The thirteen models described here differ in representation and supervision, and those two dimensions do not operate independently. A model’s representation constrains what supervision signals are meaningful, and the choice of supervision often determines which representations are practical. The next two chapters trace these interactions: first through representation (Chapter 4), then through supervision (Chapter 5).

4 Vector Representation in SVG Generation

How a model represents vector graphics internally during generation is not the same thing as what it outputs. The final result is always an SVG file. What differs is the intermediate form where shapes are created, manipulated, and decoded. The four representation types introduced in Chapter 2 each encode the same underlying geometry differently and that difference affects what the generated SVG looks like when you actually open it. This chapter works through four representation families as they appear across the chosen models, then examines what each one does to the structure of the output. Figure 4.1 shows the differences qualitatively.

4.1 Parametric Representations

Parametric representations are the purest form of how SVG already works. Every path is stored as Bézier control points, stroke attributes, and fill colors, the same parameters that appear in the file itself (Section 2.2.1). Once DiffVG made rasterization differentiable [6] (Section 2.3), these parameters became directly optimizable against image-based or semantic objectives. LIVE, SAMVG, VectorFusion, and SVGDreamer all build on this foundation [34, 35, 21, 26]. This should make parametric outputs the most editable of any representation. Every shape is an explicit geometric object a designer can select. But that breaks down at scale. On complex images, LIVE is run with up to 256 paths, added layer by layer [34]. Open that output in Illustrator or Inkscape and the problem is the image looks fine, but selecting a single object means clicking through dozens of overlapping fragments that do not correspond to any recognizable component. Fine visual detail needs more paths, not more expressive ones, so these methods are run with large path budgets. Only LIVE and SVGDreamer++ add paths during optimization [34, 38]. SVGDreamer++ [38] identifies this path entanglement explicitly as a failure mode of optimization-based parametric methods. The representation itself is not the problem. Nothing in the optimization process stops paths from piling up, which is the real issue for editability.

Which primitives do these models actually use? DiffVG can render paths, circles, ellipses,

rectangles, and polygons [6]. But most parametric models here use only one of them: closed paths made of cubic Bézier segments. LIVE and VectorFusion use four cubic segments per path [34, 21], and Im2Vec decodes every shape as a closed Bézier path [37]. A closed path can look like a circle, but it is never stored as one, so the file contains no `<circle>` or `<rect>` elements. SVGFusion is the exception. It is trained on real SVG files and keeps their circles, rectangles, and ellipses [40].

4.2 Sequential Representations

Sequential models produce some of the most compact outputs in the analyzed set. Paths are few, shapes are distinct, and the resulting SVG files are small enough to read by hand. The reason is structural. Sequential representations treat vector graphics as ordered sequences of drawing commands, mirroring the actual structure of SVG path data. DeepSVG [12] extended this formulation to structured SVG icons using a hierarchical VAE with a transformer decoder. DeepIcon [13] follows the same approach for icon vectorization. These models generate a series of commands that describe how the graphic is constructed step by step. The catch is complexity. In models that predict one command after another, such as SketchRNN (Section 2.2.2), every additional command increases the chance that prediction errors grow. A mistake appearing early in a long sequence can break everything that comes after it. DeepSVG avoids this by decoding all commands at once [12], but it is still limited by its data. It trains on SVG-Icons8 [12], which consists of simple icons with small path counts and the outputs reflect that constraint. In general most of the models that are trained on datasets have specific constraints based on dataset characteristics. DeepSVG fixes the order of commands itself and does not leave it to the model. In preprocessing, rectangles, circles, and ellipses become paths. Every icon is scaled to 256×256 . Each path then starts at its topmost-leftmost point and runs clockwise. In the ordered variant of the loss, the paths themselves are sorted lexicographically by their start position [12]. So the same icon always gives the same sequence, and the model learns one order instead of all possible ones. There is also a more fundamental limitation: a procedural command sequence does not encode explicit spatial hierarchy. The model generates commands in order but nothing in the representation directly tells you which commands belong to which semantic region.

4.3 Implicit Neural Representations

NeuralSVG [10] represents each SVG as a set of shape indices, where each index corresponds to a single shape. A small MLP learns to map these indices to Bézier control points and fill colors. So the representation is the network weights and the geometry only

exists when you query the network for it. This is fundamentally different from storing control points or commands explicitly. The model maintains stable shape correspondence during training, meaning each index always maps to the same visual region, which prevents inconsistent outputs across iterations.

NIVeL [9] follows a similar approach, encoding the entire layered structure implicitly where each layer is defined by a continuous neural function over the image domain. The outputs from both models tend to look globally coherent. Layers are clean and shapes do not bleed into each other. NeuralSVG produces compositions that hold together in ways that purely parametric outputs with hundreds of paths often do not and doing it with just 16 shapes (Figure 5.1). The cost is that individual shapes are not directly accessible until the network output is converted into discrete primitives. How clean that conversion is determines whether the final SVG is practically editable.

4.4 Latent Space as a Generative Layer

As established in Section 2.2.4, a latent space is not a geometric representation but a generative layer that sits on top of one. The geometry is always produced by a decoder that translates latent vectors back into parametric paths or sequential commands. What the latent space contributes is compression, many graphics are mapped into a shared numerical embedding learned during training. DeepSVG applies this by encoding path command sequences into latent vectors and training a transformer decoder to reconstruct them [12], building on the VAE framework (Section 2.2.4). The latent space captures structural regularities across the training dataset, so new SVGs can be generated by sampling from this space.

SVGFusion [40] extends this idea by running a diffusion process entirely within the learned latent space of a pretrained vector encoder. T2V-NPR takes a different route, arguing that global SVG-level latent spaces are too constrained for diverse path combinations and instead learns a neural path representation at the individual path level [39].

4.5 Representation Boundaries in Practice

Most models do not stay cleanly inside one representation family. VectorFusion and SVGDreamer optimize parametric primitives directly while using a latent diffusion model as the semantic prior. The parametric layer provides geometric interpretability, while the latent model provides the semantic coherence that pure geometric optimization cannot achieve alone.

NeuralSVG imposes an order on its shapes with a dropout-based regularization: during training, the shape sequence is randomly cut off, so the first shapes have to carry the coarse structure of the scene on their own [10]. Together with the fixed index mapping (Section 4.3), this makes individual shapes separable. LayerTracer [14] crosses a different boundary: its training data is sequential (step-by-step construction sequences built from about 20,000 layered SVGs made by designers), but the output is a layered parametric SVG where each layer corresponds to a meaningful visual component. The model learns layer separation as a structural convention from how designers actually build graphics. These combinations exist because no single representation handles visual realism and structural clarity at the same time.

4.6 Representation, Editability, and Structural Outcomes

Open any SVG in Illustrator or Figma and you expect to click on a shape and move it. That is the whole point. Whether an AI-generated SVG actually supports that depends more on how the model represents geometry internally than on how the rendered output looks.

The four rockets in Figure 4.1 show what this means concretely. DeepSVG works with monochrome icons like this rocket, with around 15 to 30 paths. The rocket in the figure is a sample from the SVG-Icons8 dataset, not a generated output, but it shows the kind of output DeepSVG is trained to produce. It stays at the level of an icon, not a photorealistic scene, because it is trained on simple icons with few paths (Section 4.2). A designer could open that file and rearrange things. SVGDreamer generates the same subject using 256 to 512 optimized Bézier paths [26]. Visually rich but structurally opaque, most of those paths overlap without corresponding to any recognizable component. NeuralSVG renders its rocket with exactly 16 shapes, each mapped through an MLP [10]. In the figure these read as a few large regions: a dark body with a lighter stripe, angled fins, a pale flame at the tail, and the flat background. Compared with SVGDreamer, 16 shapes are roughly a 30x reduction in element count. SVGFusion does not sit in between. Its rocket is built from only a handful of filled shapes, so in element count it is closer to DeepSVG than to SVGDreamer. What differs is the vocabulary: SVGFusion can output Bézier paths but also circles, rectangles, and ellipses [40], because its latent space was trained on real SVG files that use these elements. Table 4.1 summarizes these counts.

The number of shapes a model produces matters more for practical editability. Parametric representations should be the most editable since every shape is an explicit geometric object. But that advantage vanishes when the path count gets into the hundreds. Sequential and latent representations produce compact output within bounded domains. Implicit representations give the cleanest decomposition but at a computational cost.

DeepSVG	SVGDreamer
	
<i>sequential representation</i>	<i>parametric representation</i>
SVGFusion	NeuralSVG
	
<i>latent representation</i>	<i>implicit repr. MLP weights</i>

Figure 4.1: Rockets illustrating the three geometric representations and the latent generative layer. Top left: a rocket icon from the SVG-Icons8 dataset that DeepSVG (sequential) is trained on [12]; this is a dataset sample, not a DeepSVG output. The other three are generated outputs. Top right: SVG-Dreamer (parametric). Bottom left: SVGFusion (latent space as a generative layer; the SVG it decodes is parametric). Bottom right: NeuralSVG (implicit neural). The SVGDreamer and SVGFusion outputs are reproduced from [40], the NeuralSVG output from [10].

Table 4.1: Structural consequences of representation choice. The latent row is a generative layer on top of a geometric representation, not a geometric representation itself. Path counts reflect default configurations reported in the respective papers; n.r. = not reported.

Representation	Path Counts	Editability	Typical Output
Parametric	64–512	Low	Visually rich, paths overlap heavily
Sequential	15–30	Medium	Clean icons, breaks on complex scenes
Implicit neural	16 (NeuralSVG)	High	Few shapes, each semantically meaningful
Latent	up to 5 layers (NIVeL) Few (n.r.)	Medium	Uses full SVG vocab, looks designed

What actually drives clean structure is not the representation itself but what constraints are placed on it. Methods that use segmentation masks to initialize primitives (SAMVG) [35], saliency maps to place shapes (NeuralSVG) [10], or construction-sequence training data (LayerTracer) [14] all produce cleaner boundaries regardless of their underlying representation type. Representation defines what kinds of structures are possible. Initialization and supervision determine whether useful structure actually shows up in the output.

4.7 Summary

Representation sets the boundaries for what a model can produce. What it actually produces depends on supervision. The next chapter looks at how different supervision signals push the geometry in different directions.

5 Supervision and Guidance Strategies for SVG Generation

Representation defines how vector graphics are encoded during generation. Supervision determines what happens to those representations during training or optimization. In SVG generation, the supervision signal is whatever tells the model whether its current output is getting closer to the goal or further away. That signal might come from a target image, a text prompt scored by CLIP, a diffusion model’s learned prior, or the structure of a training dataset itself.

What makes supervision in this domain unusual is that it interacts with representation. A similar loss applied to explicit Bézier paths and to implicit neural fields produces different geometric outcomes. The representation is not the only difference between such models, but it is the biggest one (Section 7.1). This is why Chapters 4 and 5 need to be read together, neither dimension explains the structural properties of the output on its own.

Table 5.1 provides a complete overview of which signal each model uses and when it is applied.

5.1 Reconstruction-Based Supervision

In short: render the SVG, compare it to a target image, update the parameters. That is the entire loop. Reconstruction-based supervision is the most direct application of the losses described in Section 2.4.1, and for the optimization methods and Im2Vec it works through DiffVG [6]. Because the target is known, gradients provide clear information about what should change. When datasets contain clean SVG examples, models trained this way can absorb typical path counts from the data. DeepSVG, DeepIcon, and Im2Vec all use reconstruction objectives [12, 13, 37], each with a different combination of losses (Table 5.1). DeepSVG and DeepIcon compare their predictions against ground-truth SVG commands. Im2Vec has no SVG targets at all: it is trained on raster images, and its reconstruction loss compares the rendered output with the input image through

Table 5.1: Supervision signals used by each of the thirteen models analyzed in this thesis.

Model	Supervision Signal	Applied
LIVE	UDF pixel loss + Xing loss	Inference
SAMVG	MSE + LPIPS + SAM segmentation init	Inference
Im2Vec	L2 via differentiable rendering	Training
VectorFusion	SDS from Stable Diffusion	Inference
SVGDreamer	VPSD from Stable Diffusion	Inference
SVGDreamer++	VPSD + HIVE masks + reward model	Inference
NeuralSVG	SDS + fine-tuned LoRA adapter	Inference
NIVeL	SDS from DeepFloyd	Inference
DeepSVG	Cross-entropy + L1 + KL divergence	Training
DeepIcon	Cross-entropy + CLIP embedding	Training
SVGFusion	KL + latent-space denoising (VS-DiT)	Training
T2V-NPR	VSD + neural path VAE	Both
LayerTracer	Denoising on layered construction sequences	Training

differentiable rendering. Once trained, generation is a single forward pass with no further optimization (Chapter 3).

The structural quality of the output comes from the dataset, not from the loss. The loss only measures visual similarity. A model trained on messy SVGs with redundant paths would learn to reproduce that mess. A model can achieve low reconstruction error while producing overlapping geometry, because pixel-level matching does not care about what the SVG file looks like on the inside. Reconstruction supervision also requires a reference image, so it cannot support text-to-SVG generation on its own.

5.2 Semantic Guidance with Vision–Language Models

CLIP-based guidance (Section 2.4.2) can score how well a rendered SVG matches a text prompt, when no reference image exists. This makes open-ended generation possible without paired SVG datasets. CLIP measures alignment in embedding space, not structural correctness. Optimization can introduce extra primitives to push the embedding score higher, and there is nothing in the signal that penalizes fragmented output. In practice, CLIP guidance is rarely used alone. In VectorFusion [21], CLIP only ranks raster samples from the diffusion model before they are traced, and the optimization itself runs on SDS.

5.3 Diffusion-Based Supervision

Diffusion-based supervision has become the dominant strategy for text-to-SVG generation. These methods use a pretrained diffusion model as a perceptual prior, instead of relying on reference images, the same as CLIP. The diffusion model has already learned what images should look like. The SVG parameters get adjusted until the diffusion model thinks the rendered output looks right for the given text prompt.

The mechanism is Score Distillation Sampling (SDS). Variational Score Distillation (VSD) was introduced by ProlificDreamer [25] to address the oversaturation that standard SDS tends to produce, and SVGDreamer [26] builds its VPSD on it. Both are described in Section 2.4.4. What matters for this chapter is what each variant does to the geometry.

VectorFusion and NIVeL use standard SDS [21, 9]. VectorFusion takes it from Stable Diffusion. NIVeL takes it from DeepFloyd, which denoises in pixel space, so no image encoder sits in the loop [9]. NeuralSVG adds a fine-tuned LoRA adapter on top of SDS [10]. SVGDreamer replaces SDS entirely with VPSD [26]. Figure 5.1 shows the outputs side by side from the same prompt. VectorFusion produces over-smoothed geometry with 64 shapes. SVGDreamer improves diversity but uses 256 shapes to get there. NeuralSVG reaches comparable visual quality with 16 shapes.

Diffusion gradients operate in pixel space and reward perceptual realism. They do not reward geometric simplicity. Without structural constraints, these methods need a large path budget to satisfy visual objectives. Inference cost is also high, every generation requires repeated rendering and diffusion evaluation cycles, unlike dataset-driven models that produce output in a single pass.



Figure 5.1: Three diffusion-based supervision variants generating from the same prompt. VectorFusion (SDS, 64 shapes) produces over-smoothed geometry. SVGDreamer (VPSD, 256 shapes) improves diversity but uses four times as many primitives. NeuralSVG (SDS with LoRA, 16 shapes) achieves comparable visual quality with the fewest shapes. Images reproduced from [40, 10].

5.4 Dataset Supervision vs. Per-Instance Supervision

When supervision is applied changes what kind of output you get. Dataset-supervised models learn from collections of SVG graphics during training. DeepSVG [12] and DeepIcon [13] are trained on reconstruction objectives against SVG data. Im2Vec [37] is trained on raster images instead, with reconstruction through differentiable rendering. Once trained, generation requires no further optimization for any of the three. LayerTracer [14] takes this further: it trains on construction sequences derived from about 20,000 layered SVGs made by designers. Each graphic is rebuilt in a deliberate layer order that mimics how designers actually build graphics (Figure 5.2). The dataset structure itself is the supervision prior. The model learns to produce layered outputs not because a loss function rewards layering, but because every training example was built that way.




Figure 5.2: Step-by-step construction sequence from the training data of LayerTracer [14]. The frames are arranged in a serpentine grid layout, so consecutive steps sit next to each other. Each stage adds a new layer, mimicking how a designer would build the graphic incrementally. The dataset structure itself encodes the supervision prior.

Per-instance supervision works on a completely different timescale. LIVE [34], SAMVG [35], VectorFusion [21], SVGDreamer [26], NeuralSVG [10], and NIVeL [9] all optimize vector parameters from scratch for each input. No model is trained. Structure changes dynamically in response to each prompt or image, which is flexible but computationally expensive. T2V-NPR [39] combines both: it learns a path latent space from a dataset during training and then optimizes path latents for each prompt at inference. SVGFusion [40] also learns its latent space from data, but it does not optimize at inference. It samples new SVGs with its latent diffusion transformer, which takes 24 to 36 seconds per SVG.

5.5 Initialization and Structural Priors

Where optimization starts matters as much as what it optimizes for. SAMVG [35] is the clearest example. It uses the Segment Anything Model to define semantically

meaningful regions before any vector optimization begins. The loss function is still pixel reconstruction, same as LIVE. But SAMVG’s outputs have cleaner boundaries and less path fragmentation, because the starting point already encodes where objects are. The kind of supervision did not change. The initialization did, and that made the structural difference (Figure 5.3). The two outputs look almost identical at first sight, and most people would not notice a difference. But look more carefully: SAMVG’s output (right) has cleaner shape boundaries with fewer overlapping paths, while LIVE’s output (left) captures finer shading and color transitions through denser path layering. The visual fidelity is slightly higher in LIVE, but that fidelity comes from path fragmentation that makes the file harder to edit.




Figure 5.3: Both models use pixel reconstruction loss via differentiable rasterization. LIVE adds each new path where the reconstruction error is largest and produces a visually closer match to the target through dense overlapping paths, but the resulting file is fragmented and difficult to edit. SAMVG initializes from segmentation masks and produces cleaner object boundaries with fewer paths, at the cost of reduced visual fidelity [35].

T2V-NPR starts optimization inside a latent space shaped by prior training, so even a random starting latent decodes to a regular path. SVGFusion does not optimize at all, its structure comes from the training data through sampling. LayerTracer takes the idea further, its training data encodes construction order (Figure 5.2), so the model learns to produce layered outputs without any explicit structural loss. The dataset itself is the structural prior. When meaningful priors are absent, diffusion optimization needs many primitives. When structured initialization is present, the same loss produces cleaner geometry with fewer paths. Supervision tells you what the model optimizes for. Initialization tells you where it starts. That starting point often determines whether the result is editable.

5.6 Summary

Visual quality always comes first. Clean structure, when it shows up, is a side effect. Some models get closer to editability through their training data or initialization, but none of the thirteen directly optimizes for it.

6 Evaluation of SVG Generation

Chapter 5 ended with the observation that visual quality always comes first. This chapter asks whether evaluation measures anything else. Evaluating SVG generation is harder than evaluating raster image generation. A raster image is just pixels, the only thing that matters is how it looks. An SVG is a structured artifact that has to look right (semantic alignment), be geometrically stable, and remain editable. These three requirements do not always point in the same direction, and the metrics most commonly used in the literature only measure the first one. This chapter examines how the thirteen models are evaluated, which metrics are used, where those practices fall short, and what a more complete framework would need to include.

6.1 Evaluation Metrics in Current Practice

The majority of SVG generation papers evaluate their outputs using metrics borrowed directly from raster image generation research. The most common are Fréchet Inception Distance (FID) [41], CLIP-based similarity scores [20], perceptual similarity measures such as LPIPS [42], and pixel-level reconstruction metrics such as MSE. All of these share one property: they operate on rasterized outputs. The SVG file itself, its path structure, is invisible to every one of them.

FID measures the statistical distance between distributions of generated and reference images in a feature space derived from a pretrained Inception network [43]. In SVG generation, outputs are rasterized before computing FID, which means the metric captures rendered appearance only. FID also has documented technical sensitivities: it is unstable at sample sizes below 10,000 [41], and scores shift significantly depending on the reference distribution. When models are evaluated against datasets that overlap with their training data, scores appear artificially favorable. Several papers in SVG generation still compute it on substantially smaller test sets, which makes cross-paper comparison unreliable.

CLIP score measures semantic alignment between a rendered image and a text prompt [20]. It is useful for text-to-SVG tasks where the goal is semantic correctness, but like FID it operates entirely in pixel-space. A model can achieve strong CLIP scores while producing structurally fragmented or completely uneditable SVG output.

LPIPS is a perceptual similarity metric that computes feature-based distances using a pretrained network [42]. It correlates better with human perceptual judgments than MSE and is used in vectorization tasks where reconstruction fidelity against a reference image matters. SAMVG uses LPIPS as part of its optimization loss [35], and it appears as an evaluation metric too. The limitation is the same, it measures rendered appearance, not vector structure.

Mean squared error (MSE) measures pixel-level difference between rendered output and reference image. LIVE explicitly avoids it as a training loss because it produces blurry, color-averaged outputs [34], yet several papers still use it for comparability.

Human evaluation remains the most practically meaningful signal, but also the hardest to standardize. Papers typically show qualitative examples or run small user studies asking participants to rate realism, prompt alignment, or preference. These assessments capture properties that numerical metrics miss, but they are subjective and rarely reported with enough methodological detail to understand the ratings.

Table 6.1 maps which evaluation metrics each of the thirteen models actually reports. The pattern is clear: FID and CLIP dominate text-to-SVG models, MSE and LPIPS appear in vectorization models, and structural metrics are almost entirely absent across both groups. Human evaluation is present in roughly half the papers.

Table 6.1: Evaluation signals reported across the thirteen models. ✓ = reported, – = not reported.

Model	FID	CLIP	LPIPS	MSE	User	Chamfer
LIVE	–	–	–	✓	✓	–
SAMVG	✓	–	✓	✓	–	–
Im2Vec	–	–	–	✓	–	–
VectorFusion	–	✓	–	–	✓	–
SVGDreamer	✓	✓	–	–	✓	–
SVGDreamer++	✓	✓	–	–	✓	–
NIVeL	–	✓	–	–	–	–
NeuralSVG	–	✓	–	–	–	–
DeepSVG	–	–	–	✓	✓	–
DeepIcon	–	–	–	–	–	✓
SVGFusion	✓	✓	–	–	✓	–
T2V-NPR	✓	✓	–	–	✓	–
LayerTracer	✓	✓	–	✓	–	–

Three observations follow from this table. First, no model reports layer quality as a quantitative metric, despite LayerTracer, NeuralSVG, and T2V-NPR all explicitly claiming layered outputs as a contribution. Second, T2V-NPR is the only model that evaluates geometric smoothness and it does so precisely because smoothness is central to

its architectural claim, not because it is a standard metric in the field. Third, DeepIcon is the only model that reports contour-based metrics (Chamfer Distance) [13], which measure shape overlap rather than pixel similarity. No other model uses these. This pattern suggests that models tend to report metrics that favor their specific contributions while structural properties remain unmeasured. Some benchmarks have begun to recognize this gap. SVGEEditBench [44] and SVGEEditBench V2 [45] introduce structured editing tasks and evaluate outputs using raster-based, contour-based, and description-based metrics including MSE, LPIPS, and CLIP text embeddings. SVGenius [46] proposes 18 metrics across understanding, editing, and generation dimensions. However, both benchmarks are designed primarily for LLM-based SVG processing, and neither addresses the structural properties of optimization-based or dataset-driven generation methods. Zhu et al. [47] go further and measure structure directly: they remove one element at a time and score each element by how much the render changes.

6.2 Evaluation Gaps

The metrics in Table 6.1 share a fundamental limitation: none of them require parsing the SVG file. They measure what the output looks like after rasterization, not what it is made of. This means they cannot detect redundant paths, fragmented geometry, missing hierarchy, self-intersecting curves, or uneditable structure. For vector graphics, these are not secondary properties.

6.2.1 Structural Redundancy

Optimization-based methods routinely produce outputs where many small overlapping paths approximate a region that a designer would represent with one or two shapes. LIVE and VectorFusion both exhibit this behavior [34, 21]. The SVG renders correctly but is structurally opaque. No metric in Table 6.1 would detect the difference. NeuralSVG achieves a CLIP text-image similarity of 26.94 using 16 shapes, while VectorFusion reaches 26.33 with 64 shapes and SVGDreamer 26.58 with 256 [10]. A 16x difference in path count produces near-identical scores on the metric most commonly used to evaluate these models.

6.2.2 Geometric Quality

Curved paths in optimization outputs are sometimes jagged, self-intersecting, or discontinuous under zoom. T2V-NPR explicitly measures smoothness and its ablation

study shows that removing the neural path representation drops smoothness from 0.80 to 0.53 [39]. That is a large structural difference that FID would not detect. LIVE introduces a self-intersection penalty (Xing loss) during optimization [34], but does not report intersection counts as an evaluation metric. The gap between what models optimize for geometrically and what they report is notable.

6.2.3 Editability

This is the most practically important gap. The practical requirement behind this thesis (Section 1.1) is that generated SVGs must be editable by designers in standard tools. None of the metrics in Table 6.1 measure this. Editability is occasionally demonstrated by showing that a generated element can be recolored or moved, but this is not systematically evaluated across models and no paper defines criteria for what counts as editable. Human-designed SVGs already exhibit the structural clarity that generated outputs lack. Figure 6.1 shows a typical example: an icon from SVGRepo opened in Inkscape, where every shape can be independently selected and modified. The gap between what manual design delivers and what generation models are evaluated on is part of the problem: editability is a practical requirement that the field has not yet learned to measure.

6.3 Efficiency and Reproducibility

Inference speed and reproducibility are not peripheral concerns for SVG generation. They directly determine whether a method can be used in practice. Table 6.2 summarizes code availability, hardware requirements, and inference time across all thirteen models.

The efficiency gap between optimization-based and dataset-driven methods is worth considering. LIVE, the simplest reconstruction-supervised method, requires between 2,600 and 10,000 seconds per image on an RTX 3090 depending on path count [34]. VectorFusion takes 10 to 20 minutes per SVG [21]. SVGDreamer and SVGDreamer++ require 7 to 14 minutes and at least 31GB of VRAM, placing them out of reach for standard consumer hardware [26, 38]. By contrast, LayerTracer generates a layered SVG in 27 to 34 seconds [14], and SVGFusion in 24 to 36 seconds [40]. These differences are not small. Optimization-based methods take minutes to hours because they refine parameters iteratively at inference time. Dataset-driven models do it in seconds through a single forward pass.

Hardware requirements follow the same pattern but with an important exception. Dataset-driven models are faster, but several require enterprise-grade infrastructure during training, for example, SVGFusion trains on multiple A800 GPUs, and LayerTracer on a




Figure 6.1: A human-designed SVG icon from SVGRepo, opened and decomposed in Inkscape. The graphic consists of 8 Bézier paths, each representing a distinct visual component: pole, flag body, skull, eyes, and nose. Every shape can be independently selected, moved, recolored, or deleted. This structural clarity is what current evaluation metrics do not capture.

single A100 with 80GB VRAM. Once trained, these models can run inference on less demanding hardware, but the training cost creates a barrier for independent replication of these models.

Reproducibility is a separate and equally important concern for those who want to replicate the work. Code availability was recorded for every model but was not a selection criterion (Section 1.3). Five of the thirteen models have no public code release: SAMVG, VectorFusion, DeepIcon, SVGFusion, and T2V-NPR, whose repository holds only the project page. NIVeL’s code is listed as pending, so six models cannot currently be run from official code. For these models, metrics cannot be independently tested on shared test sets. When models cannot be run, cross-model comparisons rely entirely on self-reported numbers, which may reflect favorable experimental conditions.

Table 6.2: Reproducibility and efficiency overview across the thirteen models. Inference time is per-image. Hardware as reported; (train) marks training hardware. The hybrids Im2Vec and T2V-NPR (Figure 3.1) are listed with the family they share most components with.

Model	Code	Hardware	Inference Time
<i>Optimization-Based</i>			
LIVE	Public	RTX 3090	2,600–10,000 s
SAMVG	None	RTX 3090	139–2,000 s
VectorFusion	None	Not stated	10–20 min
SVGDreamer	Public	V100 (≥ 31 GB)	7–14 min
SVGDreamer++	Public	A800 (≥ 31 GB)	~7 min
NIVeL	Pending	A100	~5 min
NeuralSVG	Public	Not stated	Unknown
Im2Vec	Public	Not stated	Unknown
<i>Dataset-Driven</i>			
DeepSVG	Public	2×1080Ti (train)	Unknown
DeepIcon	None	Not stated	Unknown
SVGFusion	None	Multi-A800 (train)	24–36 s
T2V-NPR	None (project page only)	RTX 3090	~13 min
LayerTracer	Public	A100 80GB (train)	27–34 s

Table 6.2 divides along family lines: optimization-based methods are slow and generally reproducible, while several dataset-driven models are fast but require strong hardware in training step.

6.4 Toward a Multi-Dimensional Evaluation Framework

The quantitative metrics tell you how the output looks. But they tell you nothing about what the SVG file actually contains. Every metric in Table 6.1 operates on rendered pixels, not on the SVG itself. What follows is an attempt to fix that: six evaluation dimensions that cover both visual and structural quality.

6.4.1 Semantic Alignment

Does the generated SVG represent the intended concept? FID and CLIP score approximate this for text-to-SVG tasks and remain the most established signals. They should be computed on sample sizes above 10,000 and against reference distributions that do not overlap with training data.

6.4.2 Path Economy

How efficiently does the model represent visual content? A practical measure is path count relative to visually distinct regions. NeuralSVG produces semantically accurate results with substantially fewer shapes than VectorFusion or SVGDreamer for equivalent prompts [10]. This structural difference is invisible to FID. Path count is directly computable from the SVG file without rasterization.

6.4.3 Geometric Quality

Are the paths smooth and free of self-intersections? T2V-NPR directly measures curve smoothness, and its ablation study confirms that removing the neural path representation drops the smoothness score from 0.80 to 0.53 [39]. Self-intersection count is equally computable from path data. Both are SVG-native metrics that do not require rasterization and should be standard reported signals.

6.4.4 Layer Organization

Do generated layers correspond to semantically meaningful groups? This is the hardest dimension to quantify. A proxy is the ratio of semantically grouped elements to total path count. LayerTracer and NeuralSVG make the strongest architectural claims in this

dimension, but neither reports a quantitative layer quality score [14, 10]. Developing a standardized layer metric remains an open problem.

6.4.5 Editability

Can a designer modify the output using standard tools? This requires human evaluation, importing the SVG into Illustrator or Figma and assessing whether elements can be selected, recolored, and repositioned independently. Figure 6.1 shows what target editability looks like in practice. Based on their architectural design, SVGFusion, LayerTracer, and NeuralSVG are the strongest candidates in the analyzed set, though this has not been formally tested across all models.

6.4.6 Efficiency

What hardware and time does generation require? Table 6.2 provides the data. This dimension is especially relevant for industrial deployment, where generation speed and infrastructure cost are primary constraints. For dataset-driven models, training counts too. A model that needs an A100 or several A800 GPUs to train is hard to reproduce, even if inference takes seconds.

6.4.7 Applying the Framework

Table 6.3 applies these six dimensions to all thirteen models based on their architectural properties. Most ratings reflect design characteristics rather than measured outcomes. No model scores well across all six dimensions. Optimization-based methods tend to score well on semantic alignment but poorly on layer organization and path economy, and they are mixed on geometric quality. On efficiency most of them are only partial. Dataset-driven methods are fast at inference and score better on path economy. On efficiency only DeepSVG and DeepIcon score well once training hardware counts. They vary widely on semantic alignment and editability depending on training data scope. Neural implicit methods occupy a middle ground which is strong on geometry and layers, but slow. They are the only optimization-based methods rated well on layers. SVGFusion scores well across most dimensions based on its architectural design and reported metrics, but its ratings cannot be independently verified because no public implementation is available (Table 6.2). The table maps trade-offs rather than ranking models. Chapter 7 uses these ratings together with the representation and supervision analysis to compare the models directly.

Table 6.3: Proposed multi-dimensional evaluation framework applied to the thirteen models. + = good, ~ = partial, - = bad. Ratings marked * reflect metrics reported in the original paper (Table 6.1; for T2V-NPR, its smoothness score). All other ratings are the author’s assessment based on the observations from papers and the architectural analysis in Chapters 4 and 5. Editability was not measured for any model. Efficiency takes into account the hardware listed in Table 6.2. n/a = not applicable: LIVE and SAMVG reconstruct a target image, so text alignment does not apply. [†]Single forward pass; time not reported.

Model	Semantic Align.	Path Economy	Geom. Quality	Layers	Edit.	Effic.
<i>Optimization-Based</i>						
LIVE	n/a	-	~	-	-	~
SAMVG	n/a	~	~	-	~	~
VectorFusion	+	-	-	-	-	~
SVGDreamer	+	~	~	-	~	~
SVGDreamer++	+	~	+	-	~	~
NIVeL	+	+	+	+	~	~
NeuralSVG	+	+	+	+	+	~
Im2Vec	~	~	~	-	~	+ [†]
<i>Dataset-Driven</i>						
DeepSVG	~	~	~	-	~	+ [†]
DeepIcon	~	+	+	-	~	+
SVGFusion	+	+	+	+	+	~
T2V-NPR	+	+	+	-	~	-
LayerTracer	+	+	+	+	+	~

7 Comparative Synthesis

Chapters 4 and 5 analyzed representation and supervision separately, and Chapter 6 gave a way to rate the results. This chapter discusses their combined impact on the final output. The goal is to compare the methods directly, explain why some produce clean, editable SVGs while others produce visually rich outputs that fall apart when you try to edit them. Three aspects cover the chapter: how representation and supervision interact to shape what comes out, where the realism-editability trade-off shows itself, and what conditions are required to produce usable SVG output.

7.1 How Representation and Supervision Shape Structure

As already mentioned in the previous chapters, representation and supervision do not work independently. Choices of each one depend on the other and together they determine the structure of the output. Changing the representation under the same supervision signal produces a completely different output. Also, keeping the same representation and changing the supervision signal immediately impacts the result. Table 7.1 maps these combinations across all analyzed models.

Diffusion-based supervision can be a very good example. VectorFusion and SVGDreamer both use it to optimize explicit Bézier paths [21, 26]. They are run with large path budgets, up to hundreds of primitives, because that is how unconstrained parametric paths approximate fine visual detail. However, NeuralSVG uses the same kind of signal running it through an implicit neural representation and gets comparable visual quality with 16 shapes (Figure 5.1), and fewer primitives or shapes mean a better SVG file. The representation is not the only difference. NeuralSVG also adds a LoRA adapter, saliency-based initialization and dropout-based ordering, and its shape count is set in advance [10]. But the representation is the biggest of these differences, and it changes the structural character of the output.

Reconstruction losses show the same example from a different angle. LIVE and SAMVG use the same kind of pixel loss, but SAMVG starts from segmentation masks and ends up with cleaner boundaries and fewer paths (Section 5.5). That means where optimization started also made the difference.

Table 7.1: Representation type, supervision signal, and observed structural outcomes across the analyzed papers.

Representation	Supervision	Models	Structural Outcome
Parametric	Pixel recon.	LIVE, SAMVG	High path count, fragmented structure
Parametric	SDS	VectorFusion	Semantically rich, dense primitives, not editable
Parametric	VPD	SVGDreamer, SVGDreamer++	Improved diversity, object-aware decomposition, dense
Neural implicit	SDS	NIVeL, NeuralSVG	Fixed layer count, constrained structure, moderate editability
Sequential tokens	Reconstruction	DeepSVG, DeepIcon	Clean paths within training domain
Neural path latent	VSD	T2V-NPR	Smooth, stable paths, dataset-bounded
Continuous latent	Denoising	SVGFusion	Logical construction order, layers, fast inference
Layered sequences	Denoising	LayerTracer	Designer-aligned layers, explicit construction logic, strong editability

Dataset-driven methods are constrained, so some approaches may have similar — icon-wise and simple results. Representation is fixed by design and whatever structure shows up in the output is learned from the training data. DeepSVG and DeepIcon produce clean, compact icons because the SVG-Icons8 dataset contains simple graphics with small path counts [12]. The reconstruction loss teaches them to be tidy through the data. LayerTracer pushes this further: its training data encodes step-by-step construction order, so the model picks up layering as a convention rather than discovering it through optimization [14]. SVGFusion and T2V-NPR sit in between, combining learned latent structure with semantic supervision. Their outputs are more coherent than unconstrained optimization but still bounded by what the training corpus contains [40, 39].

Representations that constrain decomposition (implicit fields, learned latent spaces, layered sequences) produce cleaner outputs than representations that do not limit primitive count, regardless of the supervision signal (Tables 4.1 and 7.1). And supervision operating within a structured representation space produces more stable geometry than supervision operating directly in pixel space. NeuralSVG with SDS inside an implicit representation and T2V-NPR with VSD inside a learned path latent are both examples of this [10, 39].

7.2 The Realism–Editability Trade-off

The most consistent pattern across models is the tension between visual realism and editability. The tension is observable in specific models. Methods driven by diffusion-based supervision achieve visually detailed outputs by refining geometry until the rendered

image aligns with learned image priors. This refinement increases structural complexity, but undermines usability at the same time. VectorFusion uses 64 paths per SVG [21]. LIVE shows the same problem without any diffusion model: it represents a complex image with up to 256 paths added layer by layer [34]. In such a file, selecting or modifying a single visual component requires unlimited patience. SVGDreamer++ explicitly acknowledges this failure, noting that optimization-based approaches produce outputs where objects are entangled across paths, making component-level editing effectively impossible [38]. However, this structural cost comes with a visual benefit, as Figure 7.1 illustrates.



Figure 7.1: The realism-editability trade-off. SVGDreamer (left) produces visually detailed output through hundreds of optimized paths that overlap and cannot be individually selected in a design tool. SVGFusion (right) generates a layered SVG where each visual component is a separate, editable element, at the cost of reduced photorealistic detail [26, 40].

SVGDreamer and SVGDreamer++ consistently produce the most visually rich and semantically detailed outputs among the analyzed methods [26, 38]. Their large path budgets, guided by strong diffusion priors, fragment the structure, and that is precisely what enables photorealistic quality.

The correct choice of models and methods depends on whether the use prioritizes visual output quality or editability. Table 6.3 shows this trade-off for the parametric methods: VectorFusion and SVGDreamer are rated + on semantic alignment but – or ~ on path economy. The neural implicit methods are the exception. Models that constrain representation tend to produce outputs where fewer primitives are used. NeuralSVG represents complex scenes using a small fixed number of shapes, improving usability despite reduced photorealistic detail [10]. But even though it is constrained to a certain number of shapes, its visual quality, as shown in Figure 5.1, does not fall behind the compared models, in terms of photorealistic quality.

7.3 Conditions for Clean and Editable SVG

The analysis in Sections 7.1 and 7.2 points to five conditions that consistently separate editable outputs from fragmented ones. Table 7.2 lists them with the models that satisfy each.

Table 7.2: Five conditions for clean and editable SVG output, derived from the representation and supervision analysis in Chapters 4 and 5.

Condition	Models	Where analyzed
Primitive count control	NeuralSVG, NIVeL, T2V-NPR, SVGDreamer, SVGDreamer++	Section 4.1, Table 4.1
Semantic decomposition	SAMVG, NeuralSVG, LayerTracer	Section 5.5
Constrained supervision	NeuralSVG, T2V-NPR	Section 5.3
Informed initialization	SAMVG, SVGDreamer, NeuralSVG	Section 5.5
Design-aware priors	LayerTracer, SVGFusion	Section 5.4

No model satisfies all five. NeuralSVG comes closest on the optimization side. LayerTracer comes closest on the dataset side. NeuralSVG is more flexible but slower. LayerTracer is faster but limited to what its training data covers.

Domain coverage is a related limitation. Dataset-driven models only generate what they were trained on: icons (DeepSVG, DeepIcon) or the SVGX corpus (SVGFusion) [12, 13, 40]. Optimization-based methods handle broader domains through diffusion priors but with the structural costs described in Section 7.1. None of the thirteen covers typography, technical illustration, or character design.

8 Conclusion

This thesis examined how SVG generation methods construct vector graphics and why their structural outcomes differ. The analysis was organized around three dimensions (representation, supervision, and evaluation) and applied across thirteen models. The goal was to provide a clear reference for anyone working on SVG generation: what current models do, how their design choices shape the output, and where the field still falls short.

8.1 Answers to the Research Questions

RQ1: What are the main properties and structures that define an editable and high-quality SVG?

An editable SVG requires a small number of semantically meaningful paths, a clear layer structure where visual components correspond to distinct groupings, and stable geometry without redundant or overlapping primitives. Visual quality and editability are separable properties: a graphically rich output can simultaneously be difficult to edit if its paths are fragmented or entangled. Chapter 2 established these properties as part of the foundational definition of SVG, and Chapter 7 identified five conditions under which current models can satisfy them.

RQ2: What methods currently exist for generating SVGs from different input modalities?

The thirteen analyzed methods fall into two broad methodological families. Optimization-based methods (LIVE, SAMVG, VectorFusion, SVGDreamer, SVGDreamer++, NIVeL, and NeuralSVG) operate by iteratively refining vector parameters to match a target image or semantic description. Dataset-driven methods (DeepSVG, DeepIcon, LayerTracer, and SVGFusion) learn generation from collections of existing SVGs. Im2Vec and T2V-NPR are hybrids: Im2Vec is trained on raster images and decodes in one pass, and T2V-NPR combines a trained path prior with per-prompt optimization. Chapter 3 described their input modalities, which span raster images, text prompts, and unconditional generation

from latent distributions. Large language model approaches were reviewed but excluded from the main analysis, as discussed in Chapter 1.

RQ3: What types of deep learning architectures are used for SVG generation?

The analyzed models fall into four architecture types: per-instance optimization through a differentiable rasterizer, small MLPs that encode the geometry implicitly, transformer encoder-decoders over SVG commands (DeepSVG, DeepIcon), and VAE or diffusion transformer combinations (SVGFusion, LayerTracer). They are described in Chapter 3 and Tables 3.1 and 3.2. The structural properties of each type’s output differ more than their visual quality, which is the observation that motivated the representation-focused analysis of this thesis.

RQ4: What kinds of structure do these models produce?

Structural outcomes vary substantially across the analyzed models and correlate with representation choice more than with visual quality. Parametric optimization methods tend to produce high path counts with fragmented geometry. Neural implicit and dataset-driven methods produce more constrained outputs with interpretable layer organization, at the cost of reduced visual complexity. Chapter 4 traced how representation shapes these structural outcomes, Chapter 5 examined how supervision signals interact with them, and Chapter 7 mapped the resulting trade-offs across all thirteen models.

RQ5: How can the quality and editability of generated SVG be evaluated?

Current evaluation practice relies primarily on metrics adapted from raster image synthesis (FID, CLIP similarity, LPIPS, and MSE) which measure visual fidelity but not structural properties. Chapter 6 proposed a multi-dimensional evaluation framework covering six dimensions: semantic alignment, path economy, geometric quality, layer organization, editability, and efficiency. The chapter also identified SVGEvalBench and SVGGenius as emerging benchmarks, while noting that both were designed for language-model-based SVG generation and do not directly apply to the optimization and dataset-driven models analyzed here. Zhu et al. [47] measure structure directly, which is the direction the framework points to.

RQ6: What is the reported and assessed quality of SVG outputs across different models?

No model in the analyzed set performs well across all six evaluation dimensions simultaneously. Optimization-based models achieve the strongest visual fidelity but produce the least editable outputs. Dataset-driven models produce structurally cleaner SVGs within constrained domains. LayerTracer on the dataset side and NeuralSVG on the optimization side come closest to balancing visual quality with structural usability. Both

fall short only on efficiency, and their editability ratings are my assessment, they were not measured. SVGFusion shows similar ratings, but they remain unverifiable due to the absence of a public implementation (Chapter 6).

8.2 Contributions

The primary contribution of this thesis is an analytical framework for understanding SVG generation in terms of representation, supervision, and their structural consequences. This framework provides a consistent basis for comparing methods that differ substantially in architecture and training strategy, and makes explicit a trade-off (between visual realism and editability) that is present across the field but seldom used to compare methods.

A secondary contribution is a multi-dimensional evaluation perspective that extends beyond pixel-level metrics to structural and practical criteria. The metrics coverage table in Chapter 6 documents what each model actually reports versus what would be needed for a complete structural assessment, and identifies the specific evaluation gaps that remain open.

The analysis also produced a set of conditions for clean and editable SVG output (primitive count control, semantic decomposition, constrained supervision, informed initialization, and design-aware priors) which represent a checklist for future model design that, to my knowledge, has not been put together in this form before.

Additionally, architectural pipeline diagrams were created to visualize how the optimization-based methods share and diverge in their generation loops (Appendix C).

8.3 Limitations

The analysis is bounded by the thirteen selected models and by what their papers report. Several models with reported strong results, including VectorFusion and DeepIcon, did not have accessible code at the time of writing, so the depth of implementation-level analysis is limited for those methods. Quantitative comparison across models was not feasible because reported metrics are heterogeneous, different models evaluate on different datasets, use different implementations of shared metrics, and report results under different experimental conditions. The framework proposed in Chapter 6 is conceptual rather than being practically applicable. The ratings in Table 6.3 are my assessment from the papers, not measurements, and the examples shown are the papers' own figures. The selection of thirteen models is not exhaustive (Section 1.3).

8.4 Future Work

Three directions follow directly from the findings.

Evaluation protocols that measure structural quality remain the most immediate gap. A benchmark combining path economy, layer coherence, and editability scores alongside visual fidelity metrics would enable direct comparison across optimization-based and dataset-driven methods for the first time. This is tractable to build using existing SVG tools.

Hybrid architectures that combine fast feed-forward generation with constrained vector refinement represent the most promising near-term direction for closing the realism-editability gap. The conditions identified in Chapter 7 suggest that structural constraints can be introduced at the representation level without necessarily sacrificing visual quality, but this has not been demonstrated in a single model.

Supervision within structured representation spaces (path-level latent constraints, layer consistency objectives, construction sequence learning) appears to regulate the fragmentation that characterizes purely raster-space optimization. SVGDreamer++’s HIVE, which supervises the optimization at the level of objects and their parts, and LayerTracer’s construction-sequence supervision both point in this direction from different starting points. Converging these two approaches is a concrete research question.

8.5 Personal Takeaways

When I started this thesis, I was thinking about SVG generation from the perspective of a designer. Over time, reading papers and articles about it made me realize that there are specific patterns used across the field in its progress. In the early research phase, I thought it would be enough to find a couple of generation models and compare them. I did not even know that there are different types of methods to generate SVGs, that each method has its own architecture, representation, and supervision strategy. It was messy and complex to understand and compare them at first. But as I started analyzing the models more closely, I noticed that some patterns kept showing up. Either one model was building on a previous one, or they were using the same core architecture but adding a different tool or loss function on top. Once I saw that, I was able to categorize them, partly inspired by how other papers in the field organize their surveys. From there, it became much easier to analyze and compare them systematically.

One thing I did take away is that if someone wants to build an SVG generator without depending on token limits or API restrictions of LLM-based services, the optimization-

based and dataset-driven methods analyzed here are the reliable path. They have their own costs, mostly in hardware, but most of the models and their code are open. I think those approaches will prove more durable in the long run than closed pipelines, and working with them directly will push the field forward in ways that relying on proprietary services cannot. Anyone going that route should pay attention to the realism-editability trade-off identified in Chapter 7, and to the five conditions for clean and editable SVG output listed there. Those were the most concrete findings of this thesis, and they are the ones I would use myself if I were building something.

I regret that I was not able to compare the models quantitatively and run them locally to see the qualitative results for myself and judge whether they are practically useful. That would have made the evaluation framework in Chapter 6 stronger, and it is the most obvious direction for extending this work.

Bibliography

- [1] Anthropic. *Claude (AI Language Model)*. <https://claude.ai>. 2026.
- [2] W3Techs. *Usage Statistics of SVG for Websites*. <https://w3techs.com/technologies/details/im-svg>. Accessed: April 2026. 2024.
- [3] J. David Eisenberg. *SVG Essentials: Producing Scalable Vector Graphics with XML*. O'Reilly, 2002.
- [4] A. Quint. “Scalable Vector Graphics”. In: *IEEE MultiMedia* 10.3 (2003), pp. 99–102.
- [5] Wikimedia Commons. *Bézier Curve*. https://commons.wikimedia.org/wiki/File:Bezier_2_big.svg. Public domain. 2024.
- [6] Tzu-Mao Li, Michal Lukáč, Michaël Gharbi, Jonathan Ragan-Kelley. “Differentiable Vector Graphics Rasterization for Editing and Learning”. In: *ACM Transactions on Graphics* 39.6 (2020), 193:1–193:15.
- [7] David Ha, Douglas Eck. “A Neural Representation of Sketch Drawings”. In: *International Conference on Learning Representations (ICLR)*. 2018.
- [8] Ben Mildenhall, Pratul P. Srinivasan, Matthew Tancik, Jonathan T. Barron, Ravi Ramamoorthi, Ren Ng. “NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis”. In: *European Conference on Computer Vision (ECCV)*. 2020.
- [9] Vikas Thamizharasan, Difan Liu, Matthew Fisher, Nanxuan Zhao, Evangelos Kalogerakis, Michal Lukac. “NIVeL: Neural Implicit Vector Layers for Text-to-Vector Generation”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2024, pp. 4589–4597.
- [10] Sagi Polaczek, Yuval Alaluf, Elad Richardson, Yael Vinker, Daniel Cohen-Or. “NeuralSVG: An Implicit Representation for Text-to-Vector Generation”. In: *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. 2025, pp. 15458–15468.
- [11] Diederik P. Kingma, Max Welling. “Auto-Encoding Variational Bayes”. In: *International Conference on Learning Representations (ICLR)*. 2014.
- [12] Alexandre Carlier, Martin Danelljan, Alexandre Alahi, Radu Timofte. “DeepSVG: A Hierarchical Generative Network for Vector Graphics Animation”. In: *Advances in Neural Information Processing Systems*. Vol. 33. 2020, pp. 16351–16361.

- [13] Qi Bing, Chaoyi Zhang, Weidong Cai. “DeepIcon: A Hierarchical Network for Layer-wise Icon Vectorization”. In: *International Conference on Digital Image Computing: Techniques and Applications (DICTA)*. 2024, pp. 70–77.
- [14] Yiren Song, Danze Chen, Mike Zheng Shou. “LayerTracer: Cognitive-Aligned Layered SVG Synthesis via Diffusion Transformer”. In: *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. 2025.
- [15] Xi Liu, Chaoyi Zhou, Nanxuan Zhao, Siyu Huang. “Bézier Splatting for Fast and Differentiable Vector Graphics Rendering”. In: *Advances in Neural Information Processing Systems*. Vol. 38. 2025.
- [16] Ian Goodfellow, Yoshua Bengio, Aaron Courville. *Deep Learning*. MIT Press, 2016.
- [17] Justin Johnson, Alexandre Alahi, Li Fei-Fei. “Perceptual Losses for Real-Time Style Transfer and Super-Resolution”. In: *European Conference on Computer Vision (ECCV)*. 2016, pp. 694–711.
- [18] Karen Simonyan, Andrew Zisserman. “Very Deep Convolutional Networks for Large-Scale Image Recognition”. In: *International Conference on Learning Representations (ICLR)*. 2015.
- [19] Mark Everingham, Luc Van Gool, Christopher K. I. Williams, John Winn, Andrew Zisserman. “The PASCAL Visual Object Classes (VOC) Challenge”. In: *International Journal of Computer Vision* 88.2 (2010), pp. 303–338.
- [20] Alec Radford et al. “Learning Transferable Visual Models From Natural Language Supervision”. In: *Proceedings of the 38th International Conference on Machine Learning (ICML)*. Vol. 139. PMLR. 2021, pp. 8748–8763.
- [21] Ajay Jain, Amber Xie, Pieter Abbeel. “VectorFusion: Text-to-SVG by Abstracting Pixel-Based Diffusion Models”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2023, pp. 1911–1920.
- [22] Jonathan Ho, Ajay Jain, Pieter Abbeel. “Denoising Diffusion Probabilistic Models”. In: *Advances in Neural Information Processing Systems*. Vol. 33. 2020, pp. 6840–6851.
- [23] Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, Björn Ommer. “High-Resolution Image Synthesis with Latent Diffusion Models”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2022, pp. 10684–10695.
- [24] Ben Poole, Ajay Jain, Jonathan T. Barron, Ben Mildenhall. “DreamFusion: Text-to-3D Using 2D Diffusion”. In: *International Conference on Learning Representations (ICLR)*. 2023.
- [25] Zhengyi Wang, Cheng Lu, Yikai Wang, Fan Bao, Chongxuan Li, Hang Su, Jun Zhu. “ProlificDreamer: High-Fidelity and Diverse Text-to-3D Generation with Variational Score Distillation”. In: *Advances in Neural Information Processing Systems*. Vol. 36. 2023, pp. 8406–8441.

- [26] Ximing Xing, Haitao Zhou, Chuang Wang, Jing Zhang, Dong Xu, Qian Yu. “SVG-Dreamer: Text Guided SVG Generation with Diffusion Model”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2024, pp. 4546–4555.
- [27] Leo Sampaio Ferraz Ribeiro, Tu Bui, John Collomosse, Moacir Ponti. “Sketchformer: Transformer-based Representation for Sketched Structure”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2020, pp. 14153–14162.
- [28] Yaroslav Ganin, Tejas Kulkarni, Igor Babuschkin, S. M. Ali Eslami, Oriol Vinyals. “Synthesizing Programs for Images Using Reinforced Adversarial Learning”. In: *Proceedings of the 35th International Conference on Machine Learning (ICML)*. Vol. 80. PMLR. 2018, pp. 1666–1675.
- [29] Raphael Gontijo Lopes, David Ha, Douglas Eck, Jonathon Shlens. “A Learned Representation for Scalable Vector Graphics”. In: *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. 2019, pp. 7930–7939.
- [30] Juan A. Rodriguez, Abhay Puri, Shubham Agarwal, Issam H. Laradji, Pau Rodriguez, Sai Rajeswar, David Vazquez, Christopher Pal, Marco Pedersoli. “StarVec: Generating Scalable Vector Graphics Code from Images and Text”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2025, pp. 16175–16186.
- [31] Yiying Yang, Wei Cheng, Sijin Chen, Xianfang Zeng, Fukun Yin, Jiaxu Zhang, Liao Wang, Gang Yu, Xingjun Ma, Yu-Gang Jiang. “OmniSVG: A Unified Scalable Vector Graphics Generation Model”. In: *Advances in Neural Information Processing Systems*. Vol. 38. 2025.
- [32] Ximing Xing, Juncheng Hu, Guotao Liang, Jing Zhang, Dong Xu, Qian Yu. “Empowering LLMs to Understand and Generate Complex Vector Graphics”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2025, pp. 19487–19497.
- [33] Haomin Wang et al. “InternSVG: Towards Unified SVG Tasks with Multimodal Large Language Models”. In: *International Conference on Learning Representations (ICLR)*. 2026.
- [34] Xu Ma, Yuqian Zhou, Xingqian Xu, Bin Sun, Valerii Filev, Nikita Orlov, Yun Fu, Humphrey Shi. “Towards Layer-wise Image Vectorization”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2022, pp. 16314–16323.
- [35] Haokun Zhu, Juang Ian Chong, Teng Hu, Ran Yi, Yu-Kun Lai, Paul L. Rosin. “SAMVG: A Multi-stage Image Vectorization Model with the Segment-Anything Model”. In: *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. 2024, pp. 4350–4354.
- [36] Alexander Kirillov et al. “Segment Anything”. In: *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. 2023, pp. 4015–4026.

- [37] Pradyumna Reddy, Michael Gharbi, Michal Lukac, Niloy J. Mitra. “Im2Vec: Synthesizing Vector Graphics without Vector Supervision”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2021, pp. 7342–7351.
- [38] Ximing Xing, Qian Yu, Chuang Wang, Haitao Zhou, Jing Zhang, Dong Xu. “SVG-Dreamer++: Advancing Editability and Diversity in Text-Guided SVG Generation”. In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 47.7 (2025), pp. 5397–5413.
- [39] Peiying Zhang, Nanxuan Zhao, Jing Liao. “Text-to-Vector Generation with Neural Path Representation”. In: *ACM Transactions on Graphics* 43.4 (2024), p. 36.
- [40] Ximing Xing, Juncheng Hu, Jing Zhang, Dong Xu, Qian Yu. *SVGFusion: Scalable Text-to-SVG Generation via Vector Space Diffusion*. arXiv preprint arXiv:2412.10437 (v2). 2025.
- [41] Martin Heusel, Hubert Ramsauer, Thomas Unterthiner, Bernhard Nessler, Sepp Hochreiter. “GANs Trained by a Two Time-Scale Update Rule Converge to a Local Nash Equilibrium”. In: *Advances in Neural Information Processing Systems*. Vol. 30. 2017.
- [42] Richard Zhang, Phillip Isola, Alexei A. Efros, Eli Shechtman, Oliver Wang. “The Unreasonable Effectiveness of Deep Features as a Perceptual Metric”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2018, pp. 586–595.
- [43] Christian Szegedy, Vincent Vanhoucke, Sergey Ioffe, Jon Shlens, Zbigniew Wojna. “Rethinking the Inception Architecture for Computer Vision”. In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 2016, pp. 2818–2826.
- [44] Kunato Nishina, Yusuke Matsui. “SVGEditBench: A Benchmark Dataset for Quantitative Assessment of LLM’s SVG Editing Capabilities”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*. 2024, pp. 8142–8147.
- [45] Kunato Nishina, Yusuke Matsui. *SVGEditBench V2: A Benchmark for Instruction-Based SVG Editing*. arXiv preprint arXiv:2502.19453. 2025.
- [46] Siqi Chen et al. “SVGenius: Benchmarking LLMs in SVG Understanding, Editing and Generation”. In: *Proceedings of the 33rd ACM International Conference on Multimedia*. 2025.
- [47] Haonan Zhu, Adrienne Deganutti, Elad Hirsch, Purvanshi Mehta. *Structural Evaluation Metrics for SVG Generation via Leave-One-Out Analysis*. arXiv preprint arXiv:2604.08809. 2026.
- [48] Boris Malashenko, Ivan Jarsky, Valeria Efimova. “Leveraging Large Language Models for Scalable Vector Graphics Processing: A Review”. In: *Zapiski Nauchnykh Seminarov POMI* 546 (2025), pp. 59–80.

List of Figures

2.1	A minimal SVG file (left) and its rendered output (right). Each element is one readable line: the face, eyes, and mouth are independent objects defined by plain text parameters. Any of them can be selected and changed without touching the rest.	7
2.2	Bézier curves of degree 1, 2, and 3 (red) with their control polygons (gray). From left to right, an additional control point (blue) is added at each degree. The curve changes direction and curvature as control points are inserted or moved. Image: Wikimedia Commons [5].	8
2.3	The differentiable rendering pipeline. DiffVG converts vector parameters into a pixel image. A loss compares this image to a target. Gradients then flow back through DiffVG to update the vector parameters directly [6].	11
3.1	Taxonomy of the thirteen analyzed models by generation strategy.	16
3.2	SDS optimization loop in VectorFusion [21].	18
3.3	VPSD optimization loop in SVGDreamer [26].	18
4.1	Rockets illustrating the three geometric representations and the latent generative layer. Top left: a rocket icon from the SVG-Icons8 dataset that DeepSVG (sequential) is trained on [12]; this is a dataset sample, not a DeepSVG output. The other three are generated outputs. Top right: SVGDreamer (parametric). Bottom left: SVGFusion (latent space as a generative layer; the SVG it decodes is parametric). Bottom right: NeuralSVG (implicit neural). The SVGDreamer and SVGFusion outputs are reproduced from [40], the NeuralSVG output from [10].	27
5.1	Three diffusion-based supervision variants generating from the same prompt. VectorFusion (SDS, 64 shapes) produces over-smoothed geometry. SVGDreamer (VPSD, 256 shapes) improves diversity but uses four times as many primitives. NeuralSVG (SDS with LoRA, 16 shapes) achieves comparable visual quality with the fewest shapes. Images reproduced from [40, 10].	31
5.2	Step-by-step construction sequence from the training data of Layer-Tracer [14]. The frames are arranged in a serpentine grid layout, so consecutive steps sit next to each other. Each stage adds a new layer, mimicking how a designer would build the graphic incrementally. The dataset structure itself encodes the supervision prior.	32

5.3	Both models use pixel reconstruction loss via differentiable rasterization. LIVE adds each new path where the reconstruction error is largest and produces a visually closer match to the target through dense overlapping paths, but the resulting file is fragmented and difficult to edit. SAMVG initializes from segmentation masks and produces cleaner object boundaries with fewer paths, at the cost of reduced visual fidelity [35].	33
6.1	A human-designed SVG icon from SVGRepo, opened and decomposed in Inkscape. The graphic consists of 8 Bézier paths, each representing a distinct visual component: pole, flag body, skull, eyes, and nose. Every shape can be independently selected, moved, recolored, or deleted. This structural clarity is what current evaluation metrics do not capture. . . .	39
7.1	The realism-editability trade-off. SVGDreamer (left) produces visually detailed output through hundreds of optimized paths that overlap and cannot be individually selected in a design tool. SVGFusion (right) generates a layered SVG where each visual component is a separate, editable element, at the cost of reduced photorealistic detail [26, 40]. . . .	46
A.1	SVG generated by SVGMaker.io from the prompt “a penguin on ice with 16 shapes.” The rendered output (a) looks visually clean, but separating the elements (b) reveals 24 paths instead of 16, several of which do not correspond to recognizable visual components.	61
C.1	Reconstruction-supervised optimization pipelines. LIVE (red) and SAMVG (blue) share the same core loop: Initialized SVG Params $\rightarrow$ DiffVG $\rightarrow$ Rendered Image $\rightarrow$ Pixel Reconstruction Loss $\rightarrow$ Target Image, with gradients flowing back to update the SVG parameters. The main difference is where optimization starts. LIVE initializes directly from the target image through residual analysis. SAMVG adds the Segment Anything Model as a preprocessing step, producing semantically meaningful initialization regions before the same optimization loop begins. Same kind of loss, different starting point, different structural outcome.	68
C.2	Semantic-supervised optimization pipelines. VectorFusion (orange), SVGDreamer (orange with blue SIVE initialization), and SVGDreamer++ (purple, adding hierarchical image vectorization via Grounded SAM). All three share the same core optimization loop (Initialized SVG Params $\rightarrow$ DiffVG $\rightarrow$ Rendered Image $\rightarrow$ SDS/VPSD loss) but differ in how primitives are initialized and what additional supervision signals are applied. SVGDreamer adds semantic attention maps for initialization. SVGDreamer++ replaces those with Grounded SAM for finer part-level decomposition. . .	69

C.3	Neural implicit optimization pipelines. NeuralSVG (blue) and NIVeL (red) both store geometry in MLP weights rather than explicit Bézier coordinates, but their pipelines differ. NeuralSVG generates an initial layout through saliency-based sampling from a diffusion model, encodes shape indices via positional encoding, and optimizes MLP weights using SDS loss with a LoRA adapter on Stable Diffusion. NIVeL takes spatial coordinates as input, applies frequency encoding, and optimizes layered implicit fields using SDS with entropy loss through DeepFloyd. Both produce SVG output by converting the learned implicit representation into discrete vector primitives.	70
-----	---	----

Appendix A

Exclusion of Large Language Model Approaches

Large language models have become an active direction in SVG generation. Methods like StarVector, OmniSVG, LLM4SVG, and InternSVG treat SVG as structured code and generate it through transformer architectures pretrained on large-scale text and image collections. Commercial services followed the same path: SVGMaker.io generates SVG graphics from text prompts using GPT, Qwen, Gemini, and other models without exposing the underlying architecture or training details.

SVG is sequential XML, so language models are a natural fit for the format. The visual quality of recent LLM-based outputs has improved a lot compared to earlier sequence models. But these approaches are excluded from the main analysis for one specific reason: this thesis asks how architectural decisions in representation and supervision shape the structure of generated SVGs. In optimization-based and dataset-driven methods, that relationship can be traced directly. A parametric representation produces a certain kind of geometry. A diffusion-based loss pushes optimization in a specific direction. The connection between design choice and structural outcome is visible.

In LLM-based generation, that connection is distributed across billions of parameters. The model generates SVG tokens, but why it chose a particular path structure or how it decides on primitive count cannot be isolated the same way. This does not make LLM approaches less valuable. It makes them less suited to the kind of analysis this thesis pursues.

Figure A.1 shows what this looks like in practice. An SVG generated by SVGMaker.io from the prompt “a penguin on ice with 16 shapes” produced 24 paths in about 10 seconds. The rendered output looks clean. But open the file in a design tool and the structure tells a different story: several paths do not match any recognizable visual component, and the requested 16 shapes became 24. The output is still more editable than what some of the optimization-based models in this thesis produce, which says something

about how fast LLM-based generation is improving. But the structural decisions behind it cannot be traced to specific architectural choices, and that is exactly what this thesis needs to do.




Figure A.1: SVG generated by SVGMaker.io from the prompt “a penguin on ice with 16 shapes.” The rendered output (a) looks visually clean, but separating the elements (b) reveals 24 paths instead of 16, several of which do not correspond to recognizable visual components.

A systematic comparison of LLM-based SVG generation methods, covering both open-source and commercial systems, is a natural next step. Malashenko et al. [48] provide an initial survey of this space.

Appendix B

Model Overview Tables

The tables in this appendix collect the per-model details behind the main chapters. They follow the classification in Table 3.3 and Figure 3.1.

B.1 Task and Inference Overview

Table B.1: Task type, inference mode, input, output, supervision, and core idea of each model.

Model	Task Type	Inference	Input	Output	Supervision	Core Idea
VectorFusion	Text-to-SVG	Iterative Opt.	Text	Paths	Semantic	Distills diffusion priors via differentiable rasterization.
SVGFusion	Text-to-SVG	Forward Dec.	Text	Layers/Paths	Supervised	Continuous latent space learning using SVGX dataset.
SVGDreamer	Text-to-SVG	Iterative Opt.	Text	Paths	Semantic	Semantic-driven vectorization (SIVE) and VPSD.
SVGDreamer++	Text-to-SVG	Iterative Opt.	Text	Paths	Semantic	Hierarchical optimization with adaptive primitive control.
T2V-NPR	Text-to-SVG	Hybrid	Text	Layers/Paths	Sup. + Semantic	Dual-branch VAE learns path latent space constraints.
LayerTracer	Text-to-SVG	Hybrid	Text/Img	Layers/Paths	Supervised	Models designer construction logic via DiT sequences.
Im2Vec	Vectorization	Forward Dec.	Image	Paths	Pixel recon.	Differentiable rasterizer allows learning from appearance.
SAMVG	Vectorization	Iterative Opt.	Image	Paths	Pixel recon.	Uses SAM for semantic initialization in vectorization.
NIVeL	Text-to-SVG	Iterative Opt.	Text	Layers/Paths	Semantic	Implicit neural layers as optimization domain for SDS.
NeuralSVG	Text-to-SVG	Iterative Opt.	Text	Layers/Paths	Semantic	Scene encoded in MLP weights for layered shapes.
LIVE	Vectorization	Iterative Opt.	Image	Paths	Pixel recon.	Layer-wise hierarchical optimization of vector graphics.
DeepSVG	Unconditional	Forward Dec.	None	Paths	Supervised	Hierarchical Transformer for shape/command disentanglement.
DeepIcon	Vectorization	Forward Dec.	Image	Layers/Paths	Supervised	Hierarchical network for direct image-to-SVG translation.

B.2 Training Architecture: Optimization-Based Methods

Table B.2: Training architecture of the optimization-based methods.

Model	Trainables	Training Data	Init.	Supervision	Representation	Priors
LIVE	SVG parameters	None	Error-driven	Pixel recon. (UDF + Xing)	Parametric (flat)	None
SAMVG	SVG parameters	None	Mask-based	Pixel recon. (MSE + LPIPS)	Parametric (mask-derived)	SAM
VectorFusion	SVG parameters	None	Random / Trace	SDS (Stable Diff.)	Parametric (flat)	Stable Diff.
SVGDreamer	SVG parameters / LoRA	None	Attention- based (SIVE)	VPD	Parametric (flat)	Stable Diff.
SVGDreamer++	SVG parameters / LoRA	None	Random + HIVE	VPD + HIVE loss	Structured parametric	Stable Diff. + SAM
NeuralSVG	MLP weights + colors/ LoRA	None	Saliency- based	SDS	Implicit neural fields	Stable Diff.
NIVeL	MLP weights + layer colors	None	Random weights	SDS	Implicit neural fields	DeepFloyd
Im2Vec [†]	Encoder + RNN + 1D CNN	Raster images	Image encoding	Pixel recon. (raster only)	Parametric paths	None

[†]Im2Vec is a hybrid (Figure 3.1). It trains an encoder–decoder using differentiable rendering and reconstruction loss, but at inference there is no iterative refinement. It is placed here because its core contribution is the training objective rather than a learned generative model over SVG structure (see Section 3.1).

B.3 Training Architecture: Dataset-Driven Methods

Table B.3: Training architecture of the dataset-driven methods.

Model	Trainables	Training Data	Init.	Supervision	Representation	Priors
DeepSVG	VAE (Transformer) weights	SVG-Icons8	Latent sampling	Recon. + KL	Sequential path tokens	None
DeepIcon	Transformer weights	SVG-Icons8	CLIP embedding	Reconstruction	Sequential path tokens	CLIP encoder
SVGFusion	VP-VAE + VS-DiT	SVGX	Random latent	Recon. + Denoising	Continuous latent space	CLIP + DINO-v2
T2V-NPR [†]	NPR-VAE + Latents + LoRA	FIGR-8-SVG	Random path latents	VSD	Neural Path Repr.	Stable Diff. + CLIP
LayerTracer	DiT + LoRA	20k layered SVGs	Diffusion noise	Denoising loss	Layered pixel sequences	Flux DiT + Florence-2

[†]T2V-NPR is a hybrid (Figure 3.1). Its path VAE is trained on FIGR-8-SVG, but at inference the path latents are optimized for each prompt with VSD.

B.4 Reproducibility Overview

Table B.4: Code availability, pre-trained weights, hardware, time cost, and reproduction feasibility.

Model	Code	Pre-trained	Hardware	Cost of time	Reproduction Feasibility
LIVE	Yes	No	RTX 3090	2,600–10,000 s	Feasible on high-end consumer GPU (slow).
SAMVG	No	No	RTX 3090	139–2,000 s	Re-implementation possible but requires effort.
VectorFusion	No	No	Not specified	10–20 min	Not reproducible due to lack of public implementation.
SVGDreamer	Yes	No	V100 (≥ 31 GB VRAM)	7–14 min	Requires a high-VRAM GPU; a LoRA is trained during each run.
SVGDreamer++	Yes	No	A800 (≥ 31 GB VRAM)	~ 7 min	Requires a high-VRAM GPU; a LoRA is trained during each run.
NeuralSVG	Yes	No	Not stated	Unknown	Hardware requirements unclear.
NIVeL	Pending	No	A100	~ 5 min	High-end GPU required.
T2V-NPR	No	VAE (unknown)	RTX 3090	~ 13 min	Not reproducible (repository holds only a project page).
DeepSVG	Yes	VAE (unknown)	2 $\times$ 1080Ti (training)	Unknown	Feasible on consumer GPU.
DeepIcon	No	Transformer (unknown)	Not specified	–	Not reproducible (no official code).
SVGFusion	No	VP-VAE + VS-DiT (unknown)	Multi-A800 (training)	24–36 s	No official implementation available.
Im2Vec	Yes	VAE (unknown)	Not specified	Unknown	Likely feasible on consumer GPU.
LayerTracer	Yes	DiT + LoRA (unknown)	A100 80GB (training)	27–34 s	Enterprise GPU required.

Appendix C

Architectural Pipeline Diagrams

The following diagrams visualize the generation pipelines of selected models analyzed in this thesis. They are grouped by method family to highlight where pipelines share components and where they diverge.




Figure C.1: Reconstruction-supervised optimization pipelines. LIVE (red) and SAMVG (blue) share the same core loop: Initialized SVG Params $\rightarrow$ DiffVG $\rightarrow$ Rendered Image $\rightarrow$ Pixel Reconstruction Loss $\rightarrow$ Target Image, with gradients flowing back to update the SVG parameters. The main difference is where optimization starts. LIVE initializes directly from the target image through residual analysis. SAMVG adds the Segment Anything Model as a preprocessing step, producing semantically meaningful initialization regions before the same optimization loop begins. Same kind of loss, different starting point, different structural outcome.




Figure C.2: Semantic-supervised optimization pipelines. VectorFusion (orange), SVGDreamer (orange with blue SIVE initialization), and SVGDreamer++ (purple, adding hierarchical image vectorization via Grounded SAM). All three share the same core optimization loop (Initialized SVG Params $\rightarrow$ DiffVG $\rightarrow$ Rendered Image $\rightarrow$ SDS/VPSD loss) but differ in how primitives are initialized and what additional supervision signals are applied. SVGDreamer adds semantic attention maps for initialization. SVGDreamer++ replaces those with Grounded SAM for finer part-level decomposition.




Figure C.3: Neural implicit optimization pipelines. NeuralSVG (blue) and NIVeL (red) both store geometry in MLP weights rather than explicit Bézier coordinates, but their pipelines differ. NeuralSVG generates an initial layout through saliency-based sampling from a diffusion model, encodes shape indices via positional encoding, and optimizes MLP weights using SDS loss with a LoRA adapter on Stable Diffusion. NIVeL takes spatial coordinates as input, applies frequency encoding, and optimizes layered implicit fields using SDS with entropy loss through DeepFloyd. Both produce SVG output by converting the learned implicit representation into discrete vector primitives.

Declaration on oath

I hereby certify that I have written my master thesis independently and have not yet submitted it for examination purposes elsewhere. All sources and aids used are listed, literal and meaningful quotations have been marked as such.

April 14, 2026, Fikrat Mutallimov

Consent to plagiarism check

I hereby agree that my submitted work may be sent to PlagAware (<https://my.plagaware.com/>) in digital form for the purpose of checking for plagiarism and that it may be temporarily (max. 5 years) stored in the database maintained by PlagScan as well as personal data which are part of this work may be stored there.

Consent is voluntary. Without this consent, the plagiarism check cannot be prevented by removing all personal data and protecting the copyright requirements. Consent to the storage and use of personal data may be revoked at any time by notifying the faculty.

April 14, 2026, Fikrat Mutallimov